

World Models for Physical AI: Learning to Predict, Plan, and Act in the Real World

ANONYMOUS AUTHOR(S)

World models provide a means of representing how physical environments evolve and how an agent’s actions may change them. Their relevance to Physical AI extends across predictive representation learning, simulation, planning, policy training, and execution monitoring. These roles place different demands on what a model predicts, how actions condition its predictions, and how the predictions are used; architecture labels alone do not reveal whether a system supports prediction, decision making, or physical execution. This survey examines world models through these three questions, connecting model-based control with visual and latent prediction, structured physical representations, multimodal interaction data, and robot policies. We organize the discussion into foundations, data and model design, integration with decision making, and evaluation and deployment, and trace each system claim through its prediction target, action interface, operational role, and evaluation setting. Across the mapped evidence, three distinctions recur. First, offline predictive quality, simulated closed-loop performance, and physical closed-loop utility are related but not interchangeable. Second, system-level gains often combine data, representation, policy, planning, and inference changes, so attribution to a world-model component requires matched controls. Third, safety evidence spans conditional guarantees, empirical violation reduction, monitoring, intervention, and recovery rather than a single strength ranking. We also examine reasoning-policy hybrids as an emerging interface between general-purpose reasoning and robot control. The resulting framework connects design choices to evidence appropriate to their role while preserving the scope and denominator of each claim.

CCS Concepts: • **Computing methodologies** → **Robotics**; *Reinforcement learning*.

Additional Key Words and Phrases: world models, Physical AI, robot learning, predictive representations, vision-language-action models, model-based control

ACM Reference Format:

Anonymous Author(s). 2026. World Models for Physical AI: Learning to Predict, Plan, and Act in the Real World. *ACM Comput. Surv.* 1, 1 (October 2026), 33 pages.

1 Foundations and Scope

1.1 Introduction

Physical AI concerns agents whose decisions change a physical environment. For a robot, choosing an action requires more than identifying a goal: the system must also contend with what is currently observable, how its body can move, and how objects may respond. A world model makes some of these consequences available as predictions. The central question for this survey is how such predictions contribute to learning and acting under the constraints of physical interaction.

Predictive models can support control through different computational routes. In *World Models*, learned visual and temporal representations support policy learning, including training inside a learned environment [22]. PlaNet instead learns latent dynamics from images and uses online planning to select actions [24]. These examples motivate a distinction between using prediction to construct a policy and querying predictions while choosing an action. Their experimental settings establish particular uses of learned dynamics rather than a general guarantee of real-world competence.

2026. Manuscript submitted to ACM

Manuscript submitted to ACM

Robot policies offer a complementary interface: they map observations and instructions to actions. OpenVLA, for example, adapts a vision-language backbone using robot demonstrations to produce visuomotor behavior [41]. This interface does not by itself expose a future-state predictor that can be queried under alternative actions. The distinction is functional: a policy may contain useful knowledge of dynamics, while an explicit prediction interface provides an additional object that can be inspected, optimized against, or compared with subsequent observations.

The relationship between representation learning and action-conditioned prediction also deserves separate treatment. V-JEPA 2 combines action-free video pretraining with a subsequent action-conditioned model, V-JEPA 2-AC, for robotic planning [3]. The stages serve different purposes and should not be assigned the same control interpretation solely because they share a representation. This motivates recording both the supervision that creates a representation and the interface through which a controller later uses it.

We organize the survey around three questions: *what is predicted, how actions enter the model, and where prediction is used*. These are the survey’s comparison axes. Together they connect data and model design to operational roles without requiring all systems to share one architecture or output space. A video generator, a latent predictor, and a structured dynamics model may address different parts of this design space; their comparison should identify the relevant state variables, time horizon, and control interface before assessing performance.

The same role-based approach guides evaluation. Evidence that a model predicts observations accurately supports a different claim from evidence that its predictions improve action selection. Likewise, a closed-loop task result must be interpreted with its task distribution, controller, inference budget, and failure conditions. Our synthesis therefore records predictive performance, downstream use, and deployment conditions separately. This is an analytical choice of the survey, not a universal ranking of papers or a requirement that every supporting component demonstrate a complete robot system.

The article follows four parts. Section 1 establishes the scope and comparison axes. Section 2 examines the data, representations, and modeling choices that determine which consequences can be learned. Section 3 studies how prediction interacts with policies, planning, memory, and recovery. Section 4 organizes evaluation, reliability, safety, and open problems around the claims that a deployed system must substantiate. A companion literature catalog retains seven retrieval categories and exposes source metadata for updating and auditing the evidence base.

1.2 Working Definitions and Inclusion Boundaries

We use *world model* for a learned model that predicts aspects of an environment’s evolution from available information. For an action-conditioned world model, the inputs additionally specify actions or other interventions. The predicted variables may be observations, latent states, object relations, physical quantities, or task-relevant consequences. This working definition accommodates models learned from passive observations as well as interaction data; it does not assume that every predictive representation already provides an executable robot-action interface.

For analysis, let a history summarize the agent’s observations and past actions. A policy maps this information and a goal to an action or action sequence. An action-conditioned world model maps the history and a proposed action sequence to predicted consequences. A controller turns the selected command into realized motion. Separating these interfaces lets us ask whether a failure originates in prediction, action selection, or execution, even when some interfaces are implemented inside a shared network.

The literature corpus includes both direct methods and supporting work. Learned simulators, predictive models used for planning, and prediction-guided policy learning provide direct cases. Datasets, analytical simulators, representation encoders, sensors, and policy baselines provide supporting context according to their demonstrated roles. Their inclusion

Where does prediction enter the control loop?

Three functional interfaces for organizing the survey


A Direct action generation



B Policy proposals reviewed by a reasoner



C Action selection through explicit consequence prediction



Comparison axes Predicted variable Action interface Operational use

Conceptual interfaces. Reasoning alone does not establish an explicit future-state model.

Fig. 1. Interfaces linking observation to physical action. The schematic distinguishes direct action generation, policy proposals reviewed by a reasoning model, and candidate actions evaluated through an explicit consequence predictor. Arrows denote interface-level information flow, not the internal architecture of any one model. Observation feedback closes each loop.

does not imply that each item is itself a learned world model. In particular, the presence of a reasoning model in a control loop is insufficient to identify an explicit future-prediction mechanism.


1.3 Review Protocol and Evidence Handling

This article uses a structured narrative review with a critical evidence map; it does not claim database-complete systematic recall. The base corpus was frozen on 1 October 2026 from 929 records in seven topic categories. Fifteen rows are non-publication gap placeholders; removing them leaves 914 publication candidates and 812 normalized candidate titles. Exact identifiers, normalized titles, and logged manual decisions currently resolve these records into 758 provisional publication families. Cross-category records are retained for navigation and are not counted as independent papers. Three public world-model bibliographies were then frozen by repository commit and compared by arXiv identifier, normalized title, and manually reviewed near-title matches. Targeted searches amend this base when manuscript claims expose a missing mechanism, domain, or evaluation endpoint.

Screening separates direct world-model evidence from supporting resources and boundary comparators. A direct item predicts environment evolution or action consequences and exposes where that prediction enters learning or decision making. Datasets, policies, simulators, safety mechanisms, and reasoning systems are included only for a stated supporting role; they are not relabeled as world models. Formal publications are preferred when available, while recent

From model design to a supportable Physical AI claim

Read left to right; retain the interface, setting, and denominator



Trace, not hierarchy: a farther-right setting changes the supported claim but does not rank every paper.

Fig. 2. Claim-to-evidence trace used throughout the review. A method is described by its prediction target, relation to action, operational use, and evaluation setting before a claim is assigned. The lower checks address attribution and reporting; they are not requirements that every supporting resource satisfy a complete physical pipeline.

preprints are retained with their version and publication status. Project pages and code repositories identify resources but do not count as independent experimental validation.

For paper-level verification, we record publication identity, predicted quantity, action interface, system role, experimental setting, denominator, result, author conclusion, claim supported by this review, limitations, and source location. Ten focused evidence batches currently contain 45 such records; the manuscript cites 76 sources overall. A companion search log records exact queries where contemporaneous notes exist and marks unrecorded queries or yields as NR rather than reconstructing them retrospectively. This provides traceability for bounded, source-specific claims but not a PRISMA-style completeness guarantee: formal-version resolution, unified database searches, exclusion counts, and forward and backward citation-search logs remain incomplete.

Figure 2 turns this protocol into a claim-to-evidence trace. Its columns are questions rather than a maturity ladder: the evaluation setting changes what can be concluded, while attribution still depends on matched data, computation, candidate actions, controllers, and reporting denominators. A physical success rate can support a system-level observation without identifying which predictive component caused it.

1.4 From Learned Dynamics to Predictive Robot Systems

World-model research comprises interacting directions rather than a single replacement sequence. One direction learns dynamics for online decisions. PETS uses probabilistic ensembles and trajectory sampling to propagate uncertainty when evaluating actions [12]; PlaNet brings planning into a learned latent space [24]. Another direction uses predicted

experience to train a policy. Dreamer optimizes behavior through imagined latent trajectories [23], whereas MBPO augments policy learning with short model rollouts initialized from collected experience [34]. The distinction concerns when the model is queried, not whether the final policy is a neural network.

Physical deployment adds sensing, actuation, and interaction costs to this design space. DayDreamer applies world-model learning directly to physical platforms for manipulation, locomotion, and navigation [88]. Foundation-model approaches enter from another direction: they acquire broad visual or semantic priors before adaptation to an action interface. Recent systems such as DreamZero couple future video with action generation [94]. These developments make it useful to retain both the classical model-use distinction and the newer questions of pretraining, embodiment, and execution.

With observation–action history h_t , goal g , and candidate action sequence $a_{t:t+H-1}$, a predictive model can be written as

$$p_\theta(y_{t+1:t+H} \mid h_t, a_{t:t+H-1}, g), \quad (1)$$

where y may represent images, latent features, geometry, or other task-relevant variables. A policy instead produces $\pi_\phi(a_{t:t+H-1} \mid h_t, g)$. These expressions describe interfaces rather than a required architecture. A joint model may implement both, while a model trained on passive video may omit the action argument. The practical question is which interface is trained, queried, and tested.

1.5 Relation to Recent Surveys and Review Scope

Recent surveys already cover substantial parts of this landscape. Kirchner et al. [42] provide a control-centric review of world models for Physical AI, with explicit attention to uncertainty, structural priors, and closed-loop reliability. Hou et al. [27] organize robot world models around policy coupling, learned simulation, and visual predictive modeling. Lu et al. [59] emphasize world-action architectures and control applications. Jena et al. [35] focus on evaluation benchmarks and the comparison between prediction and executed behavior. Table 1 summarizes their stated emphases without interpreting an omitted entry as proof that a topic is absent from a paper.

Topic inclusion is therefore not the differentiator. This review uses the system claim as its comparison unit: what is predicted, how actions condition or consume the prediction, where the model enters learning or execution, and which evaluation setting supports the claim. This cross-layer trace separates direct world-model evidence from supporting resources and boundary comparators, and distinguishes offline prediction, simulated closed loop, and physical closed loop without treating them as a universal quality ranking. The seven collection categories supply literature to four narrative parts: foundations establish interfaces; data and representations determine what can be predicted; integration explains how predictions affect behavior; evaluation examines whether that use is reliable.

The present manuscript therefore makes bounded synthesis claims rather than completeness claims. Its companion protocol records the frozen catalog, provisional publication-family layer, external-library commits, deduplication order, inclusion boundaries, focused search batches, and corpus amendments. The review’s claims are bounded by that logged corpus and source-level evidence: primary sources support method descriptions, while surveys guide discovery and positioning.

Table 1. Positioning against selected 2026 surveys. The comparison concerns organizing emphasis, not topic absence or completeness.

Survey	Primary organizing emphasis	Analytical connection in this review
Kirchner et al. [42]	Physical AI: uncertainty, structure, decision coupling, and closed-loop control	Re-index design choices by prediction target and action interface, then trace them to the supporting evaluation setting.
Hou et al. [27]	Robot learning: policy coupling, simulation, and predictive models	Separate where prediction enters learning or execution and whether downstream evidence isolates its contribution.
Lu et al. [59]	World-action modeling, architecture, and control applications	Compare joint and cascaded interfaces with explicit rollout planning and reasoner-assisted action revision.
Jena et al. [35]	Benchmark protocols and predictive embodied evaluation	Attach protocol, denominator, and execution setting to concrete method claims across the four narrative parts.

2 Data, Representations, and Predictive Models

2.1 Data: From Observed Motion to Executed Actions

Training data determine which relationships a world model can learn directly and which require additional alignment. Passive video provides observations of interaction but does not necessarily contain the robot commands that produced it. Robot trajectories can pair images with commands and proprioception, while simulator rollouts can expose additional state variables. These resources provide different supervision; volume alone does not establish suitability for a particular control interface.

DROID supplies distributed real-robot manipulation demonstrations [37], while Open X-Embodiment aggregates robot datasets and studies cross-embodiment policy learning [66]. They illustrate diversity across real scenes and across robot data sources. Neither resource is itself a trained world model. Their relevance is the opportunity to learn relationships among observations, actions, and outcomes, subject to synchronization, sensing, and action conventions.

PhysInOne illustrates a different route: scale through simulation rather than action collection [113]. It contains two million rendered videos from 153,810 dynamic scenes, spanning 71 physical phenomena and asset-disjoint train, validation, and test partitions. The suite exposes geometry, motion, physical properties, and text labels, but contains neither robot command trajectories nor real-world validation. Its video-generation comparisons are also mixed across models and fine-tuning methods. PhysInOne therefore supports studying visual physical prediction under controlled synthetic variation, not the claim that scale alone yields executable or real-world dynamics.

For a predictive-control comparison, we retain the action coordinate frame, command duration, control mode, observation frequency, and embodiment. Cartesian displacements, joint targets, and gripper commands are not interchangeable numerical vectors. A video of a human moving an object does not uniquely identify a command for a particular robot. Action alignment may require robot fine-tuning, inverse dynamics, or an embodiment-specific decoder. Table 2 separates resources by role; its extraction questions are not defects asserted for every resource.

2.2 Visual Prediction and Learned Simulation

Visual world models expose future observations that can be inspected, scored, or passed to another model. This is useful when an outcome is difficult to specify as a small state vector, but depicting a desired future differs from predicting the

Table 2. Example resources organized by supervision and use. Collection corpora, learned simulators, and task benchmarks remain distinct.

Resource / example	Learning signal	Role in a WM system	Extraction priority
Passive video; V-JEPA 2 [3]	Temporal visual observations	Representation before action adaptation	Which action labels are absent?
DROID [37]	Real manipulation demonstrations	Observation–action learning	Timing, sensors, and scene splits.
Open X-Embodiment [66]	Heterogeneous robot trajectories	Shared training across sources	Body, action units, and normalization.
UniSim [92]	Learned visual transitions	Interactive synthetic experience	Supported controls and rollout validity.
LIBERO / CALVIN [53, 62]	Tasks with interactive simulation	Downstream policy evaluation	Split, success test, and reset rules.

consequence of an action. UniPi generates a visual plan from a goal and extracts control actions [14]. UniSim learns interactive visual responses to commands and uses simulated experience for downstream learning [92]. One uses a future as an intermediate plan; the other uses a predictive environment as a source of interaction.

The distinction matters when an agent queries a model repeatedly. A coherent sequence can still misrepresent contact, a grasp, or the displacement induced by a command. We therefore separate conditioning signals from generated outputs. Text-conditioned motion, inferred latent actions, and actuator commands offer different control over a rollout. A single generated clip also differs from an interactive simulator whose next input depends on its own previous outputs.

Learned simulation connects data generation with policy improvement, but creates a validation dependency. A policy can exploit regularities of the learned environment. The relevant test is not only sample plausibility but whether a policy or evaluator trained with synthetic experience performs appropriately in an independently assessed environment. We treat generated experience as an additional data source whose utility must be measured, rather than an automatic substitute for physical interaction.

Interactive World Simulator provides a recent test of both uses [86]. Its action-conditioned latent dynamics support long autoregressive interaction, synthetic demonstration collection, and policy evaluation. On the reported tasks, some imitation policies trained on 100 generated demonstrations approach counterparts trained on 100 real demonstrations, while $\pi_{0.5}$ retains a substantial real-data advantage. Policy scores in the learned simulator are positively associated with physical scores over four tasks, but the comparison uses training-distribution initial states and shows a slight optimistic bias. This evidence supports a bounded surrogate role, not unrestricted replacement of physical evaluation.

Dexterous World Models (DWM) makes the conditioning interface visually explicit [38]. Given renderings of a known static 3D scene and an egocentric hand-mesh trajectory, it generates interaction-induced residual dynamics. A 144-clip benchmark, with three random generations per sample, evaluates synthetic and real static- and dynamic-camera cases. The paper also ranks candidate hand trajectories by text–video or image similarity, but this action-evaluation result is a qualitative figure without physical execution or task-success trials. DWM consequently demonstrates visual consequence comparison for human hand-motion candidates, not closed-loop robot control.

2.3 Latent and Structured Predictive Representations

Latent models trade direct visual interpretability for representations shaped by learning objectives. Dreamer learns behavior through imagined latent trajectories [23]; TD-MPC2 uses a decoder-free latent model for trajectory optimization [25]. V-JEPA 2 separates action-free representation learning from action-conditioned adaptation for planning [3]. Feng et al. instead couple active probing with structured representation learning to seek compact, task-sufficient latent states [16]. Across DMControl, Meta-World, RoboSuite, and Franka Kitchen simulations, their single-task and transfer comparisons use five seeds, report improved sample efficiency on most tested tasks with Reacher-Hard as an explicit exception, and evaluate skill, compositional, and unseen-task transfer. The study includes no physical robot experiment, so it supports data-representation co-design in benchmark control rather than embodied deployment. Together, these studies show that “latent” identifies a prediction space, not a complete training objective or control algorithm.

Structured models make selected physical entities explicit. Graph Network-based Simulators represent systems as interacting particles and learn their evolution through message passing [70]. Tesseract instead generates RGB, depth, and surface-normal sequences, reconstructs a time-varying point cloud, and uses an inverse-dynamics model to recover actions [107]. Its nine-task RL Bench study reports higher success than its video baseline on seven tasks, but does not include physical closed-loop control. This distinction matters: a geometrically structured future can aid action decoding without being a forward model queried by candidate robot actions.

Physical Object Understanding with a Physically Controllable World Model (PhyWM) learns a 7-billion-parameter autoregressive distribution over local RGB, optical-flow, camera, and pointer tokens from three million real video clips [81]. Sparse virtual *pokes* enable conditional motion queries and support object, articulated-part, and support-relation inference. On the reported point-prompted SpelkeBench and DragAMove comparisons, its physical-motion representation improves the selected segmentation metrics over the listed baselines. However, the paper explicitly treats visual flow patches as proxies for unavailable action data; the interventions are neither calibrated forces nor robot commands. The evidence supports structure discovery from predictive distributions, not an executed manipulation policy.

DexWM makes the action abstraction itself part of the representation design [19]. It predicts DINOv2 patch latents conditioned on differences in 3D finger keypoints and camera pose, and adds a hand-consistency loss so that fingertip and wrist locations remain recoverable. The model is pretrained on 829 hours of EgoDex plus DROID, then adapted with about four hours of exploratory RoboCasa data. In simulation, its reach, place, and grasp success rates over 50 trials per task are 72%, 28%, and 58%, compared with 18%, 8%, and 14% without pretraining. A Franka–Allegro grasp study succeeds in 10 of 12 trials without physical-robot finetuning. This supports human-video pretraining and fine-grained action conditioning under that pipeline, but “zero-shot” is limited to the absence of real-robot data: the model uses simulated robot adaptation, manual grasp scoring, and two-step open-loop execution in the physical test.

ParticleFormer gives a complementary action-conditioned case. It represents rigid, deformable, flexible, and granular materials as labeled point-cloud particles, predicts their displacement under end-effector motion, and places the learned dynamics inside MPPI [30]. The paper evaluates six simulated and three physical tasks, including cloth–granular and rope–granular interaction. Its physical control evidence is limited to three rollout trials, and models are trained per scene; the result therefore supports multi-material prediction and MPC within the tested scenes, not broad cross-scene generalization.

MVISTA-4D instead predicts view-consistent RGB–depth futures from a single view and fuses them into a time-varying 3D scene [83]. It recovers actions by optimizing a trajectory latent to reproduce a generated future and then

applying a residual inverse-dynamics model. On RLBench and RoboTwin, cross-view, cross-modality, and trajectory-latent ablations each reduce manipulation success. Six physical tasks provide an additional comparison with Tesseract, but the reported offline inference requires tens of seconds on an H200. MVISTA-4D is consequently evidence for geometry as an action-inference interface, rather than real-time forward planning over candidate actions.

Making geometry explicit does not remove state-estimation and control problems. A deployed system still needs to infer state from sensors, connect commands to boundary conditions, and account for unobserved quantities. Tesseract also identifies a concrete representation limit: its single-view RGB–depth–normal sequence captures only one visible surface.

Multimodal prediction raises a related issue. Vision, proprioception, touch, and force need not answer the same question. A contact event can be weakly visible but important for execution. OmniVTA predicts short-horizon visuotactile evolution, feeds predicted tactile features to a diffusion policy, and applies a 60-Hz controller to deviations between predicted and observed touch [109]. On six groups of xArm7 tasks, the full system exceeds its no-controller variant across the reported object, geometry, and perturbation settings. Its ablations support a role for tactile prediction in this system, while the single platform and sensor configuration limit wider conclusions.

ContactWorld controls the predictive architecture and planner while varying visual and tactile representations [105]. Across 12 simulated tasks, point clouds improve average planning success over the two RGB views, while tactile force fields provide the most consistent, but not universal, tactile gain. Its 900 physical trials span six tasks, two robot platforms, and three tactile systems. The controlled Panda comparison favors point clouds combined with structured force fields; the second platform also contains cases in which adding a tactile array does not improve the corresponding visual model. This supports a conditional conclusion: touch is useful when its representation exposes task-relevant contact that vision leaves ambiguous.

Recent systems use tactile prediction at different interfaces. TouchWorld predicts short-horizon tactile subgoals for a slower goal-conditioned policy and reserves recent tactile feedback for a faster residual controller [111]. Its six-task physical study improves over generalist and tactile-policy baselines in both clean and perturbed settings, but the hierarchy changes planning, prediction, and feedback together; the full gain is not attributable to the predictor alone. DexTouch-WM instead aligns human and robot tactile layouts and retargets human motion into robot action coordinates [67]. With robot pretraining fixed at five hours, adding up to 100 hours of human interaction improves held-out robot-domain prediction. Its downstream results are deliberately mixed: policy-evaluation agreement improves on average, whereas synthetic trajectories do not consistently replace real training data across policies.

VTAM adds a two-stage alternative: it first predicts future dual-camera and GelSight latents, then freezes that backbone while training an action expert with proprioceptive state and a deformation-derived virtual-force target [97]. Each evaluated model executes 80 physical trials across chip grasping, cucumber peeling, and flat or inclined wiping. The full model reports 90%, 85%, and 95% success on the three reported task groups, while a chip-task ablation gives 9/10 successes for the full system, 1/10 without virtual-force regularization, and 0/10 for vision-only or late tactile fusion. This is direct physical evidence for the complete hierarchical visuotactile design. It does not isolate prediction quality: future video and tactile rollouts are evaluated qualitatively, the ablation has ten trials per variant, and the virtual force is an uncalibrated deformation proxy on one sensor and robot.

TacPAC moves tactile prediction into the execution interval [61]. Its WAM caches the predicted contact and the planned chunk, then an asynchronous tactile expert compares each new tactile image with that cache and rewrites only the unexecuted action suffix. Across five tasks on one Flexiv arm, with 20 trials per method and task, the vision-only, prediction-only, no-cache, and full variants average 22%, 37%, 47%, and 64% success. A correction takes 30.4 ms, and 1.6%

of calls cross at least one 30-Hz control step. This is evidence that prediction and live feedback can be complementary rather than interchangeable; the single robot, tactile hardware, and task-specific resets do not establish sensor or embodiment transfer.

Audio-WM provides a narrower non-visual case [99]. A latent flow model predicts about 2.56 seconds of future spectrogram from 1.28 seconds of audio; a separate policy consumes current and predicted audio with an image to generate 16-step Cartesian velocity commands. The closed-loop Kinova water-filling experiment reports 30 successes in 30 trials. However, the predictor is not conditioned on candidate robot actions, and the physical experiment has no matched current-audio-only or no-lookahead ablation; the simulated piano comparison states an F1 gain but does not report the number in the paper text. The evidence supports modality-specific anticipation when sound carries task progress, not a general action-conditioned physical dynamics claim.

We therefore ask what each modality makes observable, which future quantities are predicted, how signals are synchronized, and whether downstream use is prediction, planning, or feedback. WorldArena 2.0 expands evaluation toward visuotactile modalities [72]; the studies above provide model-level evidence while leaving longer compositional contact, sensor transfer, and force calibration open.

Table 3. Prediction spaces and interfaces. A system may occupy several groups. The final column states a comparison question, not a ranking.

Prediction space	Representative sources	Information exposed	Comparison question
Compact or structured dynamics			
Latent trajectories	PlaNet; Dreamer; TD-MPC2 [23–25]	Features and task-related quantities	Which losses preserve decision-relevant information?
Predictive embeddings	V-JEPA 2 [3]	Feature targets and action adaptation	What changes between pretraining and planning?
Physical structure	GNS [70]	Particle states and interactions	How is state linked to sensing and actuation?
Dexterous latent state	DexWM [19]	Patch states and finger-keypoint actions	What robot adaptation remains after human-video pretraining?
Visual and joint action interfaces			
Future images	UniPi; UniSim [14, 92]	Plans or conditioned rollouts	Is the future a target or an action consequence?
RGB–depth–normal	Tesseract [107]	Time-varying point clouds	Is geometry predicted from actions or decoded into actions?
Point-cloud dynamics and 4D futures	ParticleFormer; MVISTA-4D [30, 83]	Multi-material motion or multi-view geometry	Is the model action-conditioned or used to infer actions?
Visuotactile futures	OmniVTA; ContactWorld; TouchWorld; DexTouch-WM; VTAM [67, 97, 105, 109, 111]	Contact evolution for planning, policy, or feedback	Which gain follows prediction rather than sensing alone?
Audio futures	Audio-WM [99]	Sound-based task progress	Is audio predicted from actions or extrapolated from history?
Video and actions	DreamZero [94]	Joint video and action outputs	Which representation is executable?
Actions, state, value	Cosmos Policy [40]	Actions and outcomes for selection	When does prediction enter selection?

2.4 Generalization Across Scenes and Embodiments

Generalization should specify what changes between training and deployment. An unfamiliar object, a new goal, a shifted camera, and a different robot body impose different demands. Open X-Embodiment provides a data-centered view of heterogeneous robot learning [66]; DreamZero studies video–action modeling and transfer with different embodiments [94]. These directions motivate distinguishing transferable visual structure from the mapping to executable commands.

Simulation-to-real transfer changes the observation and dynamics distribution while preserving the action interface. Wang et al. train Cosmos Policy on roughly 3,200 automatically generated demonstrations with visual domain

randomization and deploy it without real demonstrations or finetuning [87]. On one Franka, four tasks with ten trials each average 35% success. This establishes nonzero zero-shot physical execution, not a WAM-specific advantage: the real-data Diffusion Policy references use different data and architectures, the unseen-bottle result is qualitative, and video pretraining, randomization, and joint video–action prediction are not separately ablated.

Viewpoint shift is narrower and permits stronger controls. Selective cross-view consistency constrains action, proprioception, and value outputs across same-state camera pairs while leaving future images view-dependent [29]. On LIBERO-Plus, a matched pair-trained control isolates the objective: held-out orbital extrapolation improves by 12.2 points in one seed and 15.5 in a second, whereas within-envelope interpolation does not improve. The result is simulation-only and uneven—an extreme compound shift reduces one object suite by 40.4 points—but it identifies a concrete reporting rule: separate distribution matching, interpolation, and extrapolation, and state which output coordinates should transform.

Compositional change instead asks whether familiar objects, goals, and skills can be recombined. Mishra et al. measure when an action head attends to predicted future latents rather than the current frame and use this Temporal Ratio to gate inference-time guidance [64]. Across eight physical object–receptacle tasks on one bimanual YAM, 60 trials per method give 71.7% for the unguided variant and 83.3% with guidance. The method also changes feature extraction and joint training, adds up to three video passes per replan, and can amplify an incorrect future. Zero-WAM uses human videos as in-context task specifications [112]. Its three physical task families improve over a language-conditioned baseline on held-out configurations, but use one Franka embodiment, 30 trials per family, and 252 seen-task pairs for embodiment adaptation; the prompt modality also differs between methods. These studies support bounded task recombination on a fixed platform, not unrestricted cross-embodiment transfer.

Our synthesis separates representation robustness, predictive validity, and action-interface compatibility. A system may share an encoder across robots while using separate action heads, or retain a predictor while recalibrating control. Both can be useful forms of transfer. Neither should be described as unrestricted embodiment generalization without recording which data, modules, and parameters changed during adaptation.

3 Integration with Robot Policies

3.1 Planning with Predicted Consequences

The most direct use of a world model is to compare candidate actions before executing one. PETS uses uncertainty-aware dynamics within trajectory sampling [12]; PlaNet and TD-MPC2 plan in learned latent spaces [24, 25]. Their predictive models enter an online decision procedure whose components include candidate generation, rollout, outcome scoring, and the feedback that replaces predictions with observations.

A generic receding-horizon decision can be expressed as

$$a_{t:t+H-1}^* = \arg \max_{a_{t:t+H-1}} \mathbb{E}_{p_\theta} \left[\sum_{k=0}^{H-1} \gamma^k r(y_{t+k+1}, a_{t+k}) \right], \quad (2)$$

with an initial action or chunk executed before replanning. Here r scores predicted outcomes and γ discounts future reward. This is a decision template, not an assertion that all cited systems implement the same optimizer. Learned values, terminal costs, constraints, and uncertainty penalties can change the objective. A better planner, evaluator, or candidate generator can also improve task success independently of the predictor.

Rollout length trades anticipation of delayed effects against reliance on increasingly uncertain predictions. Frequent feedback limits open-loop execution but cannot make an inaccurate model harmless: the first selected action can itself

place a robot in a difficult state. Horizon, sampling budget, and control rate therefore belong in a method comparison alongside architecture. A planning result is most interpretable when predictive quality and action-selection behavior are examined under the same task conditions.

GeoWorld shows why planning also requires a domain qualifier [103]. It combines a hyperbolic predictive representation with energy-based CEM to select procedural action-label sequences between observed and goal clips. In the matched ViT-g384 video setting, its success rate exceeds V-JEPA 2 by 1.55–2.55 percentage points across the reported CrossTask and COIN horizons, while image-based gains are less uniform. These metrics compare predicted instructional-video steps with annotated sequences; they do not execute motor commands or test physical consequences. The result supports geometry-aware procedural planning on these datasets, not embodied control.

3.2 Driving and Navigation as an Interface Test

Driving studies make the model–decision interface unusually explicit. Drive-WM conditions multiview video prediction on candidate ego trajectories and scores the generated futures with map and object rewards [85]. In its nuScenes offline evaluation, selecting among three high-level commands reduces average trajectory error and collision rate relative to a random-command condition, approaching a planner given the ground-truth command. The evidence supports candidate ranking under logged-scene metrics; it is not an interactive road deployment or a general safety guarantee.

Occupancy models expose a related distinction. OccWorld jointly generates future 3D occupancy and ego trajectory tokens [108]; the trajectory is a model output rather than a candidate whose consequences are compared. Drive-OccWorld instead conditions future occupancy on trajectory proposals and scores them with agent, road, and learned-volume costs [93]. Its cost ablations connect structured predictions to trajectory selection, but evaluation remains open-loop on nuScenes and does not remove the complete predictive module while holding the planner fixed. Joint prediction and counterfactual planning should therefore remain separate categories.

Prediction can also influence a policy only during training. DriveVLA-W0 adds autoregressive or diffusion future-image objectives to a driving VLA, yet bypasses both visual-generation paths during driving inference [48]. Its action-only comparisons and data-scale study support dense predictive supervision as a representation-learning mechanism, especially at the largest reported scale. The 70-million-frame dataset is proprietary, some smaller-scale metrics do not improve monotonically, and NAVSIM is a predictive-driver-model benchmark rather than a physical closed loop.

DriveDreamer-Policy similarly trains depth, future-video, and trajectory heads through ordered query groups, but its action head can run without explicitly generating depth or video at deployment [114]. On NAVSIM v1, adding depth or video supervision to the action-only model changes PDMS from 88.0 to 88.5 or 88.9, and using both yields 89.2; the complete model reports 88.7 EPDMS on v2. This matched ablation supports complementary predictive supervision inside this architecture. It does not establish online look-ahead or physical driving: NAVSIM evaluates logged scenes with a predictive driver model, the depth targets and depth metric both use DA3 pseudo-depth, and the ablation reports neither variance nor significance.

ADriver-I provides an earlier modular version of recurrent world–action generation [36]. It predicts speed and steering from interleaved visual–action history, generates the next frames conditioned on that action, and feeds generated frames into later predictions. Its control and video metrics are evaluated separately on nuScenes and a private highway dataset; the recurrent “driving” result is qualitative and occurs only in self-generated video. It therefore establishes an interface precedent, not evidence that imagined futures improve closed-loop driving.

DriveWAM makes the generated future an explicit action input [74]. Each inference chunk first samples a future video latent and then decodes an action chunk from that latent; new observations re-anchor the next step, while selective

video-action KV memory bounds history. On a curated 1,000-clip subset of PhysicalAI-AV, the reported 4-second ADE/FDE is 0.83/2.47. Selective memory retains 0.89/2.52 while reducing the reported 300-second cache from 3.07 to 0.25 GB relative to full KV memory. These are offline trajectory and systems measurements: the test subset is author-curated, external models use different training data, and the ablation does not isolate generated-future conditioning from the complete pretrained and jointly supervised architecture.

DAWN instead alternates action denoising with an action-conditioned short latent rollout [58]. In its matched NAVSIM v1 ablation, adding a predictor changes PDMS from 82.8 to 85.2 and enabling bidirectional world-action refinement raises it to 87.9. A 2-second latent rollout reaches 87.3 versus 87.9 at 4 seconds, while reported latency rises from 331 ms without rollout to 1,068 ms at 4 seconds. This supports short reciprocal inference within the tested model, but not deployment safety: NAVSIM and nuScenes are logged-scene protocols, no real vehicle is evaluated, and DAWN’s NAVSIM v2 aggregate score is below several listed baselines.

Navigation introduces language-conditioned long horizons. The NavForesee preprint jointly predicts short-horizon geometry, long-horizon milestones, and waypoints under a hierarchical language plan [54]. Removing either prediction horizon degrades its R2R-CE simulator results, while removing the language planner causes a larger drop. This supports complementarity within that system, not attribution of the whole gain to prediction. Its RxR-CE result is less uniformly favorable, and the current evidence is simulated and preprint-only.

Visual goal navigation provides a more direct prediction-selection interface. Navigation World Models (NWM) predicts future views under continuous actions and uses CEM either to construct a trajectory or to rerank samples from an external policy [6]. Its RECON results establish that generated futures can enter an offline trajectory objective, but they are not physical closed-loop trials; the paper also reports context loss in unfamiliar scenes and difficulty modeling pedestrians. DreamerNav instead learns a DreamerV3 policy from latent dynamics while following a static global A* path [73]. Its main comparisons use one simulated warehouse, followed by preliminary demonstrations on two quadruped platforms without a reported physical-trial denominator. These works therefore connect prediction to action selection and deployment, respectively, without jointly establishing robust real-world replanning.

Recent preprints move closer to a measured feedback loop. NavThinker queries action-conditioned next-state features and predicted pedestrian trajectories inside a social-navigation policy [28]. Its Social-HM3D ablation improves success when look-ahead features are added, but its Unitree Go2 result is a qualitative scenario without repeated-trial statistics. NavWAM jointly predicts action chunks, future views, states, and goal-progress values, then executes direct action chunks and re-queries the model after new observations [4]. In 24 image-goal episodes per method across four indoor environments, the preprint reports 19 successes for NavWAM, 14 for OmniVLA, and 4 for NWM. This is closed-loop physical evidence within one small protocol. It does not isolate future-view prediction in the physical comparison: future views and values accompany the default policy output but are not required for execution, while the predictive-head ablations are offline.

3.3 Policy Learning with Imagined Experience

A world model can influence behavior without conducting a full online search at every control step. Dreamer learns an actor through imagined latent trajectories [23]. MBPO uses short synthetic rollouts branched from real experience for policy optimization [34]. These approaches move some predictive computation into learning. They differ from online planning in deployment cost, interaction requirements, and where model errors affect behavior.

DayDreamer anchors this discussion in physical sensing and actuation [88]. More generally, imagined learning involves two coupled feedback loops: interaction updates the model, and the model changes the policy that collects

Table 4. Distinct predictive interfaces in driving and navigation. Metrics and limitations refer to each paper’s own protocol and are not a cross-paper ranking.

System	Prediction–decision interface	Evidence boundary
Drive-WM [85]	Candidate-conditioned video rollout and image-based reward	Offline nuScenes planning; three command branches.
OccWorld [108]	Joint future occupancy and ego-trajectory generation	No explicit candidate-action comparison; strongest variant uses ground-truth occupancy.
Drive-OccWorld [93]	Candidate-conditioned occupancy rollout and occupancy cost	Open-loop nuScenes metrics; protocol and ego-status choices matter.
DriveVLA-W0 [48]	Future-image auxiliary objective during training	Generation bypassed at inference; proprietary largest-scale data.
DriveDreamer-Policy [114]	Ordered depth, video, and action queries with joint predictive supervision	NAVSIM predictive-driver-model evaluation; no physical vehicle; preprint.
DriveWAM [74]	Generated future video latent conditions the next action chunk; bounded video–action memory	Author-curated logged-driving subset; no interactive vehicle loop; preprint.
DAWN [58]	Action-conditioned latent rollout recursively refines the action hypothesis	Matched NAVSIM v1 ablations; no physical vehicle; preprint.
NavForesee [54]	Dual-horizon prediction inside a language-conditioned policy	Habitat simulation; arXiv preprint.
NWM [6]	Action-conditioned future views scored by CEM or policy reranking	Logged RECON trajectories; no physical closed loop.
DreamerNav [73]	Latent imagined learning with static A* global guidance	One simulated warehouse; preliminary two-platform demonstrations.
NavThinker [28]	Candidate-action look-ahead and predicted-trajectory reward	Simulation ablation; qualitative Go2 deployment; preprint.
NavWAM [4]	Joint future, value, and action prediction in a receding feedback loop	24 physical episodes per method; one platform; preprint.

subsequent interaction. Analysis should record both loops, including the reward source and whether policy-induced distribution changes are reflected in model updates.

World-VLA-Loop makes the second loop explicit [57]. It trains an action-conditioned video model with successful and near-success trajectories, predicts both future frames and rewards, performs VLA reinforcement learning inside the model, and then uses rollouts from the updated physical policy to refine the next model. The reported two-task robot study improves after two iterations, but still requires task-specific demonstrations and new physical rollouts. The result therefore supports co-adaptation under policy shift, not elimination of real interaction.

VLAW provides a formally published physical example of the same broad co-improvement pattern [21]. On one DROID platform, it starts from 25 expert demonstrations per task category, collects 50 physical rollouts per category in each of two rounds, adapts an action-conditioned video model, and generates 500 synthetic trajectories per task for policy updates. Across five tasks, the reported mean success is 86.8% after round two, compared with 75.2% for filtered behavior cloning using the same physical rollout budget; each of these evaluations uses 50 trials per task. This 11.6-point within-paper gap supports supplemental synthetic rollouts under the study’s reward filtering and task adaptation, not replacement of real interaction or cross-platform generalization.

WorldSample keeps physical interaction in the loop while expanding each real rollout with counterfactual world-model transitions [90]. Q-aware selection and uncertainty-guided scheduling determine which synthetic transitions enter online RL. On five tasks using one Galaxea arm, its table reports average success of 82% versus 56% for HIL-SERL, with 23K versus 56K online steps. On the insertion task, removing scheduling changes success from 95% to 61% and intervention rate from 24% to 43%. The preprint does not report the denominator behind each success estimate, and the complete method also changes reward labeling, sample selection, scheduling, and model adaptation. It therefore supports selective real-synthetic online learning in a fixed task distribution, not a model-only causal claim.

WHIRL uses human intervention as a predictive training target rather than only a corrective action [95]. During physical HIL-RL, a latent world model predicts next-state reward, termination, dynamics, uncertainty, and the probability that the operator would take over; only the last term supplies actor-side risk shaping. On five Franka-LEAP tasks, the full system improves held-out autonomous success by 15–30 percentage points over its matched no-world-model HIL residual-RL baseline, while the largest reported rolling step-weighted intervention reduction is from 0.113 to 0.018 on Pull Drawer. The result is bounded by one seed and one operator per task-method cell, in-region resets, and one embodiment. It supports learning a forward risk signal from interventions, not cross-operator calibration or a runtime alarm that autonomously executes recovery.

Hi-WM moves a different human loop into the learned environment [49]. An operator interrupts failure-prone model rollouts, rewinds cached states, and supplies multiple corrective branches for post-training. Across three physical tasks and two policy backbones, the preprint reports a 19.0-point average gain over a world-model closed-loop baseline under matched post-training budgets. However, it omits the physical evaluation trial count and relies on human choices inside the simulator. The intervention is training data collection in a world model, not an autonomous alarm or physical runtime takeover.

Synthetic volume is not a sufficient measure of training value. MBPO’s restricted rollouts illustrate that model use can be selective rather than maximized [34]. For foundation-scale systems, the corresponding comparison should account for generating, filtering, and training on synthetic trajectories. Real interaction, synthetic interaction, and computation remain separate resources; a reduction in one is not automatically a reduction in all three.

3.4 Memory, Failure Detection, and Recovery

Long-horizon prediction and recovery depend on different uses of history. Memory can preserve hidden scene state or task progress; a failure detector decides that execution has departed from an expected state; a recovery mechanism selects corrective behavior. These functions may share representations, but evidence for one does not establish the others.

Mem-World uses an action-conditioned, wrist-view-centered 4D memory to retrieve geometrically relevant observations during video rollout [110]. Across five manipulation tasks, its predicted policy success rates correlate more strongly with physical success than those of its Ctrl-World comparator ($r = 0.97$ versus 0.85). Synthetic post-training raises average physical success from 58% to 72% on three long-horizon tasks, but the data pipeline retains only 20–30 successful rollouts per task from 200 candidates using human annotation. The result supports persistent simulation and data generation under that protocol; it is not an online recovery test, and the model can still predict physically invalid grasps.

WorldScape Policy 2.0 instead combines recent visual context with event-level history intended to track completed subgoals [75]. In its clean-only RoboTwin ablation, success averaged across clean and randomized tests rises from 40.91% without memory to 44.67% with short-term memory, 46.25% after adding long-term event memory, and 47.89%

after adding latent subgoal reasoning. The preprint also reports 20 physical trials per task on one dual-arm platform. Those physical comparisons vary the complete system rather than memory alone and do not inject predefined failures, so they support progress-aware execution more directly than failure recovery.

Failure-Aware RL supplies a prediction–intervention interface [45]. A short-horizon latent world model predicts future constraint violations and switches from the task policy to an offline-trained recovery policy when risk exceeds a threshold. The authors compare FARL with an offline-to-online baseline over 50 training episodes on each of three Franka tasks and report fewer failure state–action pairs. This is physical evidence for anticipatory action replacement during post-training, but the mechanism usually prevents a predicted failure rather than restoring the task after failure; it also uses task, recovery, and failure demonstrations collected for each task.

Dream2Fix uses a world model at a different point: recovery-data construction rather than runtime planning [44]. It perturbs successful real demonstrations inside Ctrl-World, verifies the generated failures, and trains a VLM to output structured trajectory corrections. On one Franka setup, 50 trials per setting yield task-level recovery rates of 46% on benchmark failures and 40% on OpenVLA deployment failures, compared with 36% and 30% for Gemini-3-Flash. Its often-quoted 81.3% result is instead offline correction-label accuracy on the generated test set. The physical evidence therefore shows partial closed-loop recovery for four controlled grasp-related failure types, not general state restoration; the world model generates training data but is not the deployed recovery policy.

Post-failure correction need not use a predictive dynamics model. ViFailback monitors a VLA with an external vision-language model and returns textual or visual-symbol guidance after detecting a failure [98]. On three previously unseen physical tasks with 21 trials per task and matched initial layouts, its two correction interfaces raise average success from 50.8–52.4% to 73.0–74.6%. This formal CVPR study supports failure diagnosis and corrective supervision, not a world-model claim. Together, these examples motivate reporting memory, prevention, detection, correction, state restoration, and task resumption as separate capabilities.

3.5 Visual Plans and Joint World–Action Models

A generated future can serve as an intermediate action representation. UniPi synthesizes goal-conditioned video and extracts actions from it [14]. This separates choosing what should happen from translating that outcome into commands. It also identifies a failure boundary: a plausible plan can be unreachable by the current robot, and a correct plan can be poorly translated into motor actions. Plan quality and action grounding require separate examination.

Joint world–action models reduce this separation by learning future observations and commands together. DreamZero builds on a pretrained video diffusion model to generate video and actions [94]. Cosmos Policy encodes actions, future observations, and values within a video model’s latent generation process, and examines direct-policy and planning modes [40]. These examples illustrate the interface convergence discussed in recent surveys [27, 59], without implying that every video-based policy performs an explicit future rollout at execution time.

Motus broadens this joint interface into switchable understanding, video-generation, inverse-dynamics, action, and video–action modes [8]. On 50 RoboTwin tasks evaluated with 100 executions per task, it reports 88.66% and 87.02% average success in clean and randomized scenes; removing the video-generation expert in a smaller matched simulation ablation reduces success from 77.00% to 25.50%. Its two-platform physical study instead reports partial subtask completion and omits the evaluation denominator. The evidence supports joint pretraining and a useful video expert in the tested simulation, plus physical feasibility, but not a general performance ranking or explicit online planning.

JOPAT changes which future variables are coupled to action generation [20]. It jointly denoises actions, visual latents, 2D point tracks, and visibility, then executes action chunks in a receding-horizon loop. With data and optimization held fixed on LIBERO-Long, the joint model reports 96.4% average success versus 77.4% for latent-only prediction. Across four LeRobot SO-101 tasks with ten rollouts per method and task, the same comparison is 57.5% versus 35.0%. These matched ablations support point correspondences as a complementary control interface within JOPAT. They do not establish broad superiority: the physical sample is small, precision insertion remains at 1/10, and track supervision inherits errors from an off-the-shelf tracker.

VAMPO instead changes the training objective of the predictive branch [18]. It post-trains a video prediction model with GRPO and an expert-future latent reward before training the action-generation module. On CALVIN ABC→D, the base average completed-sequence length is 4.28; post-training only the predictor while retaining the original action module gives 4.35, and retraining the action module on the revised representation gives 4.56. The decomposition supports objective alignment as a viable post-training route, while also showing that the larger gain follows downstream adaptation rather than predictor replacement alone. The preprint reports no variance, its checkpoints fluctuate from 4.31 to 4.56, and its physical figure omits evaluation denominators, so it does not yet establish a general causal link from latent accuracy to physical control.

V-JEPA Policy provides a matched test of future-latent supervision inside a policy [102]. A frozen V-JEPA 2.1 encoder defines the visual space; a future predictor supplies layer-wise context states to a flow-matching action expert, but does not condition its predictions on candidate actions. With data, architecture, training seed, and update budget fixed, removing the future loss while retaining future queries changes success from 97.25% to 91.65% on LIBERO and from 79.25% to 68.81% on LIBERO-Plus. Predictor-only pretraining on DROID, followed by matched downstream training, changes the two physical task results from 11/20 and 7/20 to 15/20 and 17/20. These comparisons support explicit predictive supervision and transferable predictor initialization within this architecture, not explicit model-based search or a cross-paper ranking: the pretraining adds 100,000 upstream updates, the ablation uses one training seed, and the physical test covers one dual-arm platform and two tasks.

Two neighboring JEPA-style systems place prediction earlier in the pipeline. VLA-JEPA learns latent transition tokens from current observations using future V-JEPA targets, then adapts a flow-matching action head [77]. In its within-paper human-video ablation, LIBERO-Plus changes from 62.9% to 79.5%, but LIBERO changes only from 96.1% to 97.2% and the paper reports no benefit on one SimplerEnv robot average. JEPA-WAM instead uses a joint current-future target to shape a shared policy predictor and removes the transition branch at deployment [52]. Under one matched LIBERO-Plus recipe, disabling transition prediction while retaining V-JEPA features scores 77.0%, versus 79.2% for the full design. Its five-task physical study averages normalized completion over 10 rollouts per task and setting, rather than binary success. These results support training-time predictive representation shaping within the reported systems; neither method supplies candidate-action-conditioned online rollout, and neither reports multi-seed training uncertainty.

JEPA-VLA moves prediction farther upstream: it feeds frozen V-JEPA 2 embeddings from the two most recent frames into an otherwise direct policy, without a downstream future target or transition predictor [63]. In the authors' OpenVLA-OFT implementation, adding the encoder and gated fusion changes average LIBERO success from 90.30% to 96.40%, and the six-task RoboTwin easy/hard averages from 54.8/9.3% to 73.5/17.7%. This supports the complete representation-fusion package, not the predictive objective alone: encoder capacity, fusion layers, and compute all change, no multi-seed uncertainty is reported, and the physical study covers one task with an incompletely specified per-condition denominator. JEPA-VLA is therefore evidence for predictive pretraining as a visual prior, not for a downstream world model or online imagination.

NoiseGate treats the reliability of each imagined video latent as another control variable inside joint denoising [32]. It freezes the video–action backbone and uses simulator reward to train a per-latent timestep scheduler. In its 50-task RoboTwin configuration, average success changes from 92.58% for the demonstration-trained backbone to 94.28% after scheduler training. A separate 15-task, smaller-backbone ablation reports 57.5% for shared timesteps, 61.3% after independent-noise training, 63.4% with a hand-crafted schedule, and 67.5% with learned scheduling. The protocols must remain separate: the result is simulation-only, uses a single reported seed, and spends additional online rollout interaction to optimize the scheduler.

WALL-WM changes the temporal unit rather than only the objective or schedule [47]. It aligns language, predicted video, and action at semantic event boundaries and supports variable-length event execution. On internal real-robot suites, a pretrained event-mode system exceeds a pretrained fixed-length baseline without view-interaction attention in Task Progress, from 32.6 to 71.6 on reasoning manipulation and 22.0 to 53.75 on generalization. This comparison simultaneously changes cross-view attention and execution format, and the paper does not report physical trial counts or uncertainty. Together with acknowledged platform-aligned data and tuning advantages, it supports event-level execution as a promising combined system design, not a single-factor causal claim or a cross-platform result.

WAV combines a future-state model, a trajectory-value model, and action decoding through iterative latent inference [46]. Removing its latent planning stage reduces its reported LIBERO average from 98.1 to 96.4, with the largest drop on the Long suite. Its three-task physical comparison also favors WAV over a baseline that shares its video-pretrained model, although each task uses a separate model and only 15 trials. This is evidence for an implicit planning interface under those conditions, rather than a universal advantage over explicit search.

LaRA-VLA marks a useful boundary [5]. It predicts a next-observation latent from the current image, instruction, and reasoning state, then uses inverse-dynamics supervision and an action expert to produce commands. The predicted future token is not conditioned on candidate actions and is not queried to compare their consequences. On SimplerEnv, its matched ablation raises average success from 55.2% without chain-of-thought components to 68.8% with latent textual and visual reasoning; its physical study uses 12 trials for each of four long-horizon task categories. These results support predictive latent supervision inside a policy, but not an action-conditioned forward model or counterfactual planner. We therefore retain LaRA-VLA as a boundary comparator rather than a core world-model instance.

DynBridge and MM-ACT sharpen the same distinction [51, 82]. DynBridge trains interaction tokens with future-trajectory reconstruction and action imitation; its five-task Franka study uses 10 trials and compares the full policy with an internal no-dynamics ablation. MM-ACT jointly supervises text, future images, and actions: on eight RoboTwin tasks, adding text and image generation to the action-only setting raises average success from 43.13% to 52.38%, while its three-task Franka comparison uses 20 trials per task but does not repeat that component ablation. Both support prediction as policy supervision under their protocols. Neither exposes a candidate-action-conditioned forward model for counterfactual search, and MM-ACT’s physical result cannot isolate the predictive objective.

Table 5 compares the location of prediction in the pipeline. For a joint system, a predicted outcome may be consumed by action selection, shape training through an auxiliary loss, or remain another generated output. The mechanisms can coexist, but suggest different ablations. They should not be collapsed into an undifferentiated statement that a policy “uses a world model.”

WorldSync addresses whether a learned simulator follows actions outside expert demonstrations [11]. Its WorldEcho protocol pairs visual-integrity gating with end-effector trajectory alignment over expert and feasible off-expert queries. In a matched two-round policy-improvement study on one physical cup-stacking task, both policies start at 48% success; the WorldSync condition reaches 68%, while an expert-trained CtrlWorld condition reaches 56%. The paper does not

Table 5. Prediction in policy learning and execution. Suggested comparisons are review criteria, not experiments asserted for every cited paper.

Integration route	Example	Model use	Controlled comparison
Online planning	PETS; TD-MPC2 [12, 25]	Candidate evaluation	Fix candidates and budget; vary predictor.
Imagined learning	Dreamer; MBPO [23, 34]	Policy/value training	Fix real-data budget; vary synthetic use.
Visual plan to action	UniPi [14]	Plan generation and decoding	Separate plan errors from decoder errors.
Joint generation	DreamZero; JOPAT [20, 94]	Coupled video, structured motion, and actions	Vary future-state variables with matched training.
Predictive auxiliary	LaRA-VLA; DynBridge; MM-ACT [5, 51, 82]	Latent, trajectory, or image supervision	Match policy and data; remove only prediction target.
Predictive encoder fusion	JEPA-VLA [63]	Frozen history embeddings condition a direct policy	Match encoder size, fusion capacity, latency, and data.
Training-time JEPA target	VLA-JEPA; JEPA-WAM [52, 77]	Latent transition supervision; predictor removed at deployment	Match policy and data; isolate target, representation, and human-video budget.
Future-informed action expert	V-JEPA Policy [102]	Predictor context states condition action flow	Retain future queries; remove only future loss.
Predictor post-training	VAMPO [18]	Reward alignment of future latents	Hold the action module fixed before joint adaptation.
Adaptive denoising	NoiseGate [32]	Per-latent reliability scheduling inside joint generation	Match the frozen backbone; report rollout budget, seeds, and variance.
Event-grounded execution	WALL-WM [47]	Variable-length video-action events	Separate event boundaries, cross-view modules, and pretraining.
Joint model plus selection	Cosmos Policy [40]	Direct policy or planning	Direct versus planning mode; report compute.
Implicit latent planning	WAV [46]	Iterative future/value inference	Remove planning with the same model and data.
Co-evolving simulator and policy	World-VLA-Loop; VLAW [21, 57]	Model-based RL or synthetic data plus model refresh	Track physical data and gains per iteration.
Reasoner assistance	Astra studies [17, 76]	Proposal review or editing	Fix policy; vary observation and prediction inputs.

state the physical trial denominator, and its component ablations are performed in RoboTwin rather than on the robot. The evidence connects broader action coverage and intervention-effect supervision to downstream improvement in that protocol, but does not isolate a universal action-faithfulness mechanism.

3.6 Reasoning-Policy Hybrids: The Astra Case

Recent Astra evaluations motivate distinguishing direct numerical action generation from a hybrid interface that reviews a learned policy’s proposals. In Su et al.’s technical report [76], Astra can correct actions proposed by $\pi_{0.5}$. On

Table 6. Interfaces and study scope in the Astra discussion. This is a provenance table, not a cross-paper performance leaderboard.

Source	Interface examined	Scope retained in the comparison
Su et al. [76]	Direct action generation; review of $\pi_{0.5}$ proposals	Selected tasks and aligned cases; public policy scores are reference results.
Zhang et al. [101]	LLM as policy	Full 42-task RoboDojo evaluation for Astra.
Galbot Team [17]	Direct and hybrid interfaces across domains	Domain-specific protocols; overlap with the earlier report.
Black Forest Labs [9]	FLUX 3 Action proposals with Astra review	Report-specific hybrid protocol and resource accounting.

ten selected RoboDojo tasks with five aligned cases per task, the authors report 48% success for the hybrid and 26% for direct Astra control. The public-policy comparisons use published reference scores rather than same-seed reruns. These findings support a bounded discussion of complementary action-generation and correction roles; they do not isolate a future-prediction mechanism.

A separate evaluation by Zhang et al. [101] covers all 42 RoboDojo tasks, with 50 episodes per task for Astra. It reports 22.48% average success over 2,100 trials and a capability profile that varies substantially across task types. This broader result should be discussed alongside the selected-task hybrid study, while keeping their task sets and denominators separate.

The Galbot team’s September 29 preprint [17] extends the investigation across six domains, including cooperation with learned policies. Its gripper-manipulation discussion includes the same 48% RoboDojo hybrid result. We treat the overlapping result as part of one evidence series rather than a second independent replication. The paper also documents inference constraints, including paused physics during its locomotion reasoning calls. This makes the relation between simulated success and time-critical control an explicit part of the discussion.

Black Forest Labs’ FLUX 3 Action report [9] examines another combination: a fast action policy supplies proposals that Astra may accept, edit, replace, or stop. The report connects hybrid performance to inference cost and elapsed time, alongside task completion. This provides an industrial case for evaluating the entire decision loop, with the report’s own protocol and provenance retained.

For world-model research, these interfaces raise a specific attribution question. Does a hybrid improve because it predicts action consequences, identifies the correct target, repairs a proposal, or uses more computation? A useful comparison would hold the base policy, observations, control limits, and inference budget fixed while varying the information supplied to the reviewer: current observations, proposed actions, and explicit predicted consequences. This comparison is a proposed evaluation design of the survey, not an experiment established by the cited reports. It connects the Astra discussion to the article’s main question: when, and through which interface, does prediction change action selection?

4 Evaluation, Reliability, and Research Directions

4.1 Prediction Quality and Decision Utility

Evaluation should begin with the model’s claimed role. A representation learner needs evidence about relevant downstream features; an action-conditioned predictor needs tests of the relationship between commands and outcomes;

a planner needs evidence about executed behavior. These are complementary questions, not a single evidence ladder that every dataset, theoretical result, and component must climb.

WorldArena compares perceptual quality with functional uses of embodied world models, including data generation, policy evaluation, and planning [71]. WorldArena 2.0 expands toward visuotactile modalities, interactive learning, and multiple simulated and physical platforms [72]. This broadens evaluation beyond generated-video appearance. The resulting protocols provide evidence about tested functions and settings, rather than universal certification of a high-scoring model.

WorldLens makes this separation concrete for driving world generation through 24 dimensions covering generation, 4D reconstruction, action following, downstream perception, and human preference [50]. Its action-following protocol includes both non-reactive evaluation and a reactive generative simulator in which policy outputs update the ego state and condition the next generated frame. The highest reported route completion among its tested generators is only 13.51%, despite more favorable appearance scores elsewhere. This exposes a functional gap, but not road safety: the closed loop covers five simulated sequences across two maps, uses fixed pretrained planners, and includes no physical vehicle.

Interactive World Simulator illustrates why the functional role also needs an external check [86]. Besides reporting video metrics over 192-step rollouts, it trains policies on generated demonstrations and compares policy rankings in the learned and physical environments. The observed positive association is more relevant to policy selection than pixel error alone, but it is not a calibration guarantee: the study covers four correlation tasks, training-distribution initial configurations, and a simulator that is slightly optimistic on three of them.

Policy benchmarks answer a related question. LIBERO studies knowledge transfer in robot learning [53]; CALVIN evaluates language-conditioned long-horizon manipulation [62]. Both can assess a complete system, but task success alone does not identify the responsible component. Jena et al. [35] foreground this distinction. Even within the same named suite, task subsets, demonstrations, initialization, and compute must be aligned before attributing differences to prediction.

ManipArena applies this principle to physical manipulation [78]. Its 15 tabletop tasks use ten trials each—four in-distribution, four controlled shifts, and two held-out conditions—with subtask-level partial credit. Seven policy configurations produce at least 1,050 physical trials, yet the best aggregate success rate is 38%. More importantly, data sampling, language granularity, and single- versus multi-task finetuning alter rankings within the same architecture. These results support controlled attribution and expose unsaturated performance; ten trials per task, two held-out trials, embodiment-aligned pretraining, and scoped mobile evaluation preclude a general WAM-VLA ranking.

4.2 Uncertainty, Constraints, and Runtime Guardrails

Safety requires an account of what can go wrong, which quantity is constrained, and how the system responds. These choices lead to different claims. PETS propagates model uncertainty through candidate trajectories [12]; CPO optimizes reward under expected-cost constraints [1]; neither mechanism alone establishes that a physical robot will avoid every hazardous event. The uncertainty or cost must be tied to an action rule, a safety definition, and an operating distribution.

Formal certificates are conditional. Berkenkamp et al. combine a statistical dynamics model with Lyapunov verification and prove high-probability stability within an estimated region of attraction [7]. The result assumes, among other conditions, Lipschitz dynamics and policies, calibrated confidence bounds, a Lyapunov function, and an initial safe policy. Its experiment is a simulated inverted pendulum. Thomas et al. instead penalize trajectories that a short-horizon model predicts will become unsafe [80]. Their theorem requires a specified unsafe set, a finite horizon within which

Table 7. Evaluation targets and reporting requirements. These criteria match evidence to claims without assigning a universal evidence grade.

Target	Example setting	Informative observation	Protocol detail
Predictive content	Held-out trajectories	Error, consistency, or likelihood	Variable, horizon, units, and split.
Action sensitivity	Same history; varied commands	Predicted differences match observed consequences	Action semantics and intervention coverage.
Functional use	WorldArena series; Interactive World Simulator [71, 72, 86]	Utility as data source, evaluator, planner, or environment	Specific role and downstream system.
Task performance	LIBERO; CALVIN [53, 62]	Completion and sequential outcomes	Tasks, resets, trials, data, and budget.
Attribution under shift	ManipArena; stage-wise degradation [43, 78]	Where rankings or errors change	Paired tasks, training recipe, perturbation, and confidence interval.
Safety behavior	SafeDreamer; Recovery RL [31, 79]	Violations, interventions, and recovery	Cost, exposure, severity, and success.
Runtime feasibility	Physical closed-loop execution	Delay and valid control outputs	Hardware, sensing, action duration, and paused physics.

irrecoverable states fail, and a calibrated set-valued model; experiments on four MuJoCo tasks report fewer, not zero, violations.

Uncertainty-aware shielding offers a third interface. Yu et al. learn a distributional reachability certificate and a shield policy, reporting comparable return with fewer training violations on four simulated tracking and navigation tasks [96]. The formal T-ASE article contains no physical-robot experiment and lists real-system deployment as future work. SafeDreamer integrates cost-aware planning and Lagrangian updates with a learned world model [31]. Across Safety-Gymnasium simulation tasks it improves the reward–cost tradeoff, but its nearly zero-cost result is specific to the benchmark cost definition and does not certify open-world physical safety.

Recovery separates task progress from corrective behavior. Recovery RL learns task and recovery policies, using prior data to identify situations requiring intervention [79]. We include it as a complementary safety mechanism, not relabel it as a generative world model. Control barrier functions can restrict control through a specified safe-set condition [2]; their conclusions still depend on state estimation, dynamics, control authority, and feasibility.

At the system level, Kim et al. divide foundation-model-enabled robot safety into action, decision, and human-centered dimensions and propose monitoring and intervention layers [39]. This is an architectural position paper without an evaluation of the proposed stack. It therefore motivates separate decision and action gates, but does not show that the architecture reduces physical incidents. Across these mechanisms, we distinguish guarantees under stated assumptions, empirical violation reduction, recovery demonstrations, and unvalidated design proposals.

Recent systems instantiate narrower parts of this architecture. Liu et al. use latent future rollouts to request human supervision before predicted manipulation failures [55]; two physical tasks improve over a matched-intervention baseline, although each task uses one human supervisor and the paper does not report a single physical-trial denominator. RoboGuard translates rules and a semantic world model into temporal-logic constraints for an LLM planner [68]. On a physical Jackal, it blocks all tested RoboPAIR attacks across 35 harmful prompts while preserving 10 safe behaviors, but assumes an accurate world model and does not cover all dynamics-dependent risks. FEARL separates a foundation-model

controller from a small verifiable safety-action module [13]. In 18 physical navigation episodes, shielding changes the collision rate from 27.8% to zero with the same 61.1% task success; this evidence remains conditional on calibrated low-dimensional safety sensors and the verified input domain.

Failure monitoring is not yet intervention. Foresight feeds action-conditioned V-JEPA 2-AC latents to a causal detector and calibrates time-varying alarms with functional conformal prediction [100]. On teleoperated physical rollouts, its Transformer reports ROC-AUC values of 0.93, 0.87, 0.79, and 0.89 across four policy-embodiment settings, leading three settings but not the SmolVLA setting. The detector is evaluated offline: alarms do not stop the robot, request help, or trigger recovery. Its significance level is selected per benchmark for balanced accuracy, cross-policy transfer is asymmetric, and inference takes 183.64 ms on an H200, over 99% of it in the world-model backbone. Moreover, the stated false-alarm guarantee requires exchangeability between calibration and deployment data. Thus, the study supports action-conditioned failure monitoring, not reduced physical incidents under distribution shift.

DreamAvoid connects prediction to a physical action change under a narrower protocol [15]. A causal trigger invokes test-time video rollout and value-based reranking at one to three critical phases per trajectory. Across four real tasks with 40 trials per task and method, its failure-aware variant succeeds in 116/160 episodes, versus 78/160 for the base policy and 87/160 for GPC-RANK. The intervention takes about 2.133 s and pauses high-level commands while low-level servos hold position. A post-hoc audit of 44 failures finds two untimely triggers; among the other 42, six hazardous actions are scored optimistically, while 36 low-valued actions are still executed because they remain best in the candidate pool. The evidence supports selective physical reranking in slow precision tasks, not certified avoidance, abstention, or safety under distribution shift.

Integrity of the predicted future is a separate concern. Chen et al. perturb the observation-to-imagination path of three imagine-then-act models under a white-box threat model [10]. The important result is conditional: attacking a future latent barely changes a reactive policy, whereas a LaDi-WM sampling MPC that explicitly trusts imagined futures falls from 14/20 successes under matched random noise to 1/20 under the adversarial direction at one perturbation level. This is a single simulated task, and direct vision shares an encoder with imagination refinement, so it does not establish a universal WAM vulnerability. It does establish that a verifier or planner consuming prediction needs its own integrity test; policy robustness alone is not sufficient.

FoMo-FD tests a related monitor in surgery-inspired manipulation [33]. It calibrates a flow-matching world model at a fixed $\alpha = 0.05$ using 19 successful rollouts per task. The combined two-simulation, two-physical benchmark contains 320 failed and 80 held-out successful rollouts; the wrist-view detector reports 96.6% failure detection at a 1.3% false-alarm rate. Each physical failure mode has 10 rollouts and each physical task has 20 held-out successes. These results are also computed offline. A measured 13.98 Hz scoring rate exceeds the experiment’s 10 Hz control loop, but computational feasibility does not establish timely physical intervention or reduced harm.

Selective assistance is a separate decision interface. KnowNo requests help when a conformal prediction set for an LLM planner is not a singleton; in its 50-trial physical tabletop study, it matches a simple set baseline’s plan success while reducing stepwise help from 0.72 to 0.58 [69]. CoFineLLM explicitly finetunes a planner to reduce prediction-set size and help rate and reports a hardware out-of-distribution demonstration [84]; the published abstract does not by itself establish perception or low-level control calibration. Neither method is a world model. They expose a useful reporting target: coverage must be paired with how often, how long, and at what task stage a person intervenes. ThriftyDagger makes this burden explicit through intervention frequency, duration, and context switching [26]; its physical evidence is a single cable-routing task with 15 intervention-aided trials, so it is a measurement precedent rather than broad reliability evidence.

Table 8. Complementary safety mechanisms. Their scope follows their assumptions and evidence; rows are not ordered by safety strength.

Mechanism	Source	Function	Required check
Stability certificate	Berkenkamp et al. [7]	Certify a region of attraction	Initial safe set, confidence coverage, smoothness, and dimension.
Near-future avoidance	Thomas et al. [80]	Penalize predicted unsafe rollouts	Unsafe-set definition, model calibration, and sufficient horizon.
Reachability shield	Yu et al. [96]	Override actions using uncertainty-aware risk	Coverage, intervention rate, false alarms, and physical transfer.
WM cost planning	SafeDreamer [31]	Plan with learned costs and rewards	Cost-prediction error, exposure, and observed violations.
Recovery control	Recovery RL [79]	Switch to corrective behavior	Missed hazards, false alarms, and recovery success.
Predictive recovery gate	FARL [45]	Predict near-future violations and replace risky actions	Prediction calibration, threshold, demonstrations, and post-failure scope.
Modular guardrails	Kim et al. [39]	Separate monitoring and decision/action gates	Implemented modules, latency, independence, and incident reduction.
Predictive monitor	Liu et al. [55]	Request human help from latent rollouts	Trial denominator, supervisor variation, and intervention burden.
Selective future reranking	DreamAvoid [15]	Pause at critical phases and replace the nominal action chunk	Trigger misses, unsafe best candidate, hold latency, and distribution shift.
Prediction-integrity probe	Trusted-imagination attack [10]	Test a downstream consumer that relies on imagined futures	Threat model, random control, task-level effect, and shared encoder paths.
Calibrated failure monitor	Foresight [100]	Alarm from action-conditioned predicted latents	Distribution match, operating point, lead time, online response, and latency.
Semantic plan guard	RoboGuard [68]	Convert rules and a world model into plan constraints	Map integrity, rule coverage, adaptive attacks, and dynamics.
Verified safety module	FEARL [13]	Certify a small sensor-to-action module	Sensor calibration, verified domain, specification coverage, and transfer.

4.3 Operational Reliability and Attribution

Reliability depends on the complete observation–decision–execution loop. A planner can act on a stale observation; a controller can saturate; an action chunk can continue after the scene changes. This motivates reporting end-to-end latency and action timing, not just network inference time. The Astra evaluations in Section 3 also illustrate why simulation timing must be explicit: pausing an environment removes consequences of reasoning delay that remain present during continuous physical operation.

Stage-wise evaluation can localize this loop without assuming that one metric predicts the final outcome. Lee and Zhang apply ten synthetic sensing degradations to the same 50 OGBench drawer tasks and track representation shift, future-prediction residual, planner preference over a shared candidate pool, and task success [43]. The relative ordering changes across stages, and every primary success-rate confidence interval includes zero. This supports paired multi-stage diagnosis, not a claim that a degradation improves behavior; the study is simulation-only and its latest-observation restoration probe does not execute the recovered trajectory.

A broader WAM-VLA robustness comparison reaches a similarly conditional result [104]. On RoboTwin 2.0-Plus, several WAMs are strong under lighting, noise, and layout changes but weak under camera and robot-initial-state shifts; on LIBERO-Plus, a VLA achieves the highest aggregate result. Training data, finetuning, model size, and denoising steps differ, and some LIBERO numbers come from prior reports. The evidence therefore maps perturbation profiles and runtime costs rather than establishing a causal advantage for either model family.

The same caution applies when spending more inference compute. GeoBoN samples multiple WAM rollouts and selects an action chunk by cross-view depth-reprojection consistency [106]. Across five simulated benchmark-backbone settings and four seeds, its $N = 8$ selector improves aggregate success by roughly 1–2 percentage points; a gate triggers extra sampling at 26.2% of decision points on average and retains 74.8% of the always-on gain. The gate still adds latency, some task groups decline, and larger candidate pools can select false low-score outliers. This is evidence for selective prediction-based reranking under calibrated multi-view geometry, not a free or physically validated scaling law.

Attribution requires controls at the relevant interface. To assess a predictive planner, hold candidate generation and scoring fixed where possible. To assess predictive pretraining, compare training signals with matched data and compute. For a reasoning-policy hybrid, distinguish additional observations, additional inference, proposal editing, and explicit consequence prediction. Not every study needs every comparison; the comparisons determine which explanation the study supports.

A practical reporting unit is an episode with its outcome, elapsed time, intervention history, and failure type. Aggregation should retain denominators and variation across tasks, seeds, or environments. Reports that reuse experiments should be linked as evidence families when appropriate. This is important for rapidly released systems where an early project report and an expanded preprint may contain overlapping trials.

Recovery also needs a stated endpoint. LIBERO-Recover separates action retry, action adaptation, task-relevant object-state restoration, and environmental restoration across more than 1,000 simulated scenarios [56]. Six evaluated policies perform substantially better on the two action-level categories than on state-level recovery, and mixed recovery training improves recovery-set results without consistently improving standard LIBERO success. Because the benchmark is simulated and starts from constructed failure states, it does not by itself establish online failure recognition or physical recovery. It does provide a useful reporting distinction: avoiding a predicted hazard, correcting one action, restoring the state, and resuming the original task are different outcomes.

REBOOT operationalizes the endpoint for real precision assembly data [65]. Its 2,160 episodes across 18 tasks include equal numbers of expert and human recovery demonstrations, with timestamps for failure onset, recovery initiation, and return to a valid continuation state. The failure distribution comes from three learned policies, but a single operator reproduces the states and teleoperates every correction; no autonomous recovery policy is evaluated. The resource therefore supports phase-level data and measurement claims, not deployed recovery performance.

Two non-world-model systems further expose the interface boundary. VLA-FixBench evaluates a pause-diagnose-rollback-correct loop on two Franka tasks [91]. Off-the-shelf VLM assistance attains 0–25% success versus 30% for the VLA-only baseline, because false interventions can corrupt nominal executions; the reported 35-point physical

improvement is a human-expert upper bound, not an autonomous result, and the paper does not state the physical trial denominator. AgentChord instead precompiles likely failure branches and rejoins the nominal task graph after correction [89]. Across six physical tasks with 20 trials each and one or two injected disturbances per trial, it reports 77.5% average task success versus 72.5% for the strongest listed comparator. This is direct evidence of task resumption under anticipated faults on one platform, but it depends on predefined monitors and recovery branches and is not evidence for a world model. Together, these studies sharpen the remaining target: calibrated online detection, corrective control, physical state restoration, and final task completion under one protocol, with cross-platform and cross-operator replication.

4.4 Research Directions and Synthesis

From prediction scores to useful decisions. An important question is how changes in predictive quality alter decisions made with the model. This requires examining variables used by a planner, not just a visual aggregate score. Functional benchmarks provide one entry point [71]. Further comparisons should examine cases in which plausible predictions misrank actions or omit consequential alternatives.

A recent risk-informed world-model perspective sharpens this gap as three mismatches: likelihood versus risk, prediction versus intervention, and finite prediction horizon versus accumulated consequences [60]. Its proposed representation, counterfactual reasoning, memory, and assurance components form a research agenda rather than an evaluated architecture. The testable question is whether preserving rare, action-changing evidence alters action ranking, intervention, recovery, or abstention under matched information and compute.

Contact, hidden state, and embodiment. Structured and multimodal models raise a concrete systems question: which hidden variables must be sensed or inferred for contact prediction, deformable interaction, touch and force, and robots with different control conventions? Comparative studies should distinguish gains from observing a variable, predicting it more accurately, and changing how a policy uses it. This distinction prevents sensor access, model quality, and control benefit from being conflated.

Adaptation and computation. Online updates, longer memory, and larger inference budgets change a system’s operating conditions. Comparisons should report adaptation data, regression, compute, and behavior under delayed or failed inference. Prediction and control need not run at the same rate. The design question is how to allocate computation while maintaining a response when prediction is uncertain, unavailable, or too slow.

Synthesis. World models connect experience to action through learned simulation, predictive representation, online planning, imagined learning, and joint world–action generation. These routes differ in what is predicted and where predictions are consumed. The review follows those dependencies from data and representation to behavior and reliability. In the evidence mapped here, studies of contact-rich interaction, navigation and driving, and safety under distribution shift use different platforms, endpoints, and often narrow closed-loop protocols. These differences preclude a uniform performance comparison. A defensible synthesis must therefore preserve each source’s prediction target, action interface, evaluation setting, and denominator rather than flatten heterogeneous protocols into one trajectory of progress.

References

- [1] Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. 2017. Constrained Policy Optimization. In *Proceedings of the 34th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 70)*. PMLR, 22–31. <https://proceedings.mlr.press/v70/achiam17a.html>

- [2] Aaron D. Ames, Samuel Coogan, Magnus Egerstedt, Gennaro Notomista, Koushil Sreenath, and Paulo Tabuada. 2019. Control Barrier Functions: Theory and Applications. In *2019 18th European Control Conference*. IEEE, 3420–3431. doi:10.23919/ECC.2019.8796030
- [3] Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zhohov, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, Francisco Massa, Marc Szafraniec, Kapil Krishnakumar, Yong Li, Xiaodong Ma, Sarath Chandar, Franziska Meier, Yann LeCun, Michael Rabbat, and Nicolas Ballas. 2025. V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning. arXiv:2506.09985 <https://arxiv.org/abs/2506.09985v1> Preprint, version 1, 11 June 2025.
- [4] Daichi Azuma, Taiki Miyanishi, Koya Sakamoto, Shuhei Kurita, Yaonan Zhu, Petr Khrapchenkov, Motoaki Kawanabe, Yusuke Iwasawa, and Yutaka Matsuo. 2026. NavWAM: A Navigation World Action Model for Goal-Conditioned Visual Navigation. arXiv:2606.13494 <https://arxiv.org/abs/2606.13494v1> Preprint, version 1, 11 June 2026.
- [5] Shuanghao Bai, Jing Lyu, Wanqi Zhou, Zhe Li, Dakai Wang, Lei Xing, Xiaoguang Zhao, Pengwei Wang, Zhongyuan Wang, Cheng Chi, Badong Chen, and Shanghang Zhang. 2026. Latent Reasoning VLA: Latent Thinking and Prediction for Vision-Language-Action Models. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*. PMLR, 5409–5426. <https://proceedings.mlr.press/v306/bai26h.html>
- [6] Amir Bar, Gaoyue Zhou, Danny Tran, Trevor Darrell, and Yann LeCun. 2025. Navigation World Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 15791–15801. https://openaccess.thecvf.com/content/CVPR2025/html/Bar_Navigation_World_Models_CVPR_2025_paper.html
- [7] Felix Berkenkamp, Matteo Turchetta, Angela P. Schoellig, and Andreas Krause. 2017. Safe Model-based Reinforcement Learning with Stability Guarantees. In *Advances in Neural Information Processing Systems*, Vol. 30. <https://proceedings.neurips.cc/paper/2017/hash/766ebcd59621e305170616ba3d3dac32-Abstract.html>
- [8] Hongzhe Bi, Hengkai Tan, Shenghao Xie, Zeyuan Wang, Shuhe Huang, Haitian Liu, Ruowen Zhao, Yao Feng, Chendong Xiang, Yinze Rong, Hongyan Zhao, Hanyu Liu, Zhizhong Su, Lei Ma, Hang Su, and Jun Zhu. 2026. Motus: A Unified Latent Action World Model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 35101–35113. https://openaccess.thecvf.com/content/CVPR2026/html/Bi_Motus_A_Unified_Latent_Action_World_Model_CVPR_2026_paper.html
- [9] Black Forest Labs. 2026. FLUX 3 Action: A 7B World Action Model for Robot Control. Technical report. <https://bfl.ai/models/flux-3-action> Accessed 1 October 2026.
- [10] Linghan Chen, Kaiyan Ji, and Minyu Guo. 2026. Attacking the Trusted Imagination: Oracle-Level Integrity Attacks on Imagine-then-Act World Models. arXiv:2606.22966 [cs.LG] <https://arxiv.org/abs/2606.22966v1> Preprint, version 1, submitted 22 June 2026.
- [11] Sixiang Chen, Jiaming Liu, Jixian Wu, Yichen Guo, Tinghao Wang, Siyuan Qian, Hao Chen, Jiajun Cao, Jian Tang, and Shanghang Zhang. 2026. Do Robotic World Models Really Follow Actions? Diagnosing and Aligning Action-Conditioned Generation for Policy Learning. arXiv:2608.24885 <https://arxiv.org/abs/2608.24885v1> Preprint, version 1, 25 August 2026.
- [12] Kurtland Chua, Roberto Calandra, Rowan McAllister, and Sergey Levine. 2018. Deep Reinforcement Learning in a Handful of Trials using Probabilistic Dynamics Models. In *Advances in Neural Information Processing Systems*, Vol. 31. <https://proceedings.neurips.cc/paper/2018/hash/3de568f8597b94bda53149c7d7f5958c-Abstract.html>
- [13] Davide Corsi, Kyungmin Kim, and Roy Fox. 2026. Verifiable Foundation Models for Robot Safety. arXiv:2606.23754 <https://arxiv.org/abs/2606.23754v1> Preprint, version 1, 22 June 2026.
- [14] Yilun Du, Sherry Yang, Bo Dai, Hanjun Dai, Ofir Nachum, Josh Tenenbaum, Dale Schuurmans, and Pieter Abbeel. 2023. Learning Universal Policies via Text-Guided Video Generation. In *Advances in Neural Information Processing Systems*, Vol. 36. https://proceedings.neurips.cc/paper_files/paper/2023/hash/1d5b9233ad716a43be5c0d3023cb82d0-Abstract-Conference.html
- [15] Xianzhe Fan, Yuxiang Lu, Shenyuan Gao, Xiaoyang Wu, Ruihua Han, Manling Li, and Hengshuang Zhao. 2026. DreamAvoid: Critical-Phase Test-Time Dreaming to Avoid Failures in VLA Policies. arXiv:2605.11750 [cs.RO] <https://arxiv.org/abs/2605.11750v2> Preprint, version 2, revised 23 September 2026.
- [16] Fan Feng, Yujia Zheng, Minghao Fu, Yongqiang Chen, Guangyi Chen, Kevin Patrick Murphy, Biwei Huang, and Kun Zhang. 2026. Learning Task-Sufficient World Models by Synergizing Agentic Exploration and Structured Modeling. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*. PMLR, 30616–30641. <https://proceedings.mlr.press/v306/feng26aa.html>
- [17] Galbot Team. 2026. Systematically Exploring the Capabilities of GPT-6 Astra as Embodied Policies. arXiv:2609.38537 <https://arxiv.org/abs/2609.38537v1> Version 1, 29 September 2026.
- [18] Zirui Ge, Pengxiang Ding, Baohua Yin, Qishen Wang, Zhiyong Xie, Yemin Wang, Jinbo Wang, Hengtao Li, Runze Suo, Wenxuan Song, Han Zhao, Shangke Lyu, Zhaoxin Fan, Haoang Li, Ran Cheng, Cheng Chi, Huibin Ge, Yaozhi Luo, and Donglin Wang. 2026. VAMPO: Policy Optimization for Improving Visual Dynamics in Video Action Models. arXiv:2603.19370 <https://arxiv.org/abs/2603.19370v1> Preprint, version 1, 19 March 2026.
- [19] Raktim Gautam Goswami, Amir Bar, David Fan, Tsung-Yen Yang, Gaoyue Zhou, Prashanth Krishnamurthy, Michael Rabbat, Farshad Khorrami, and Yann LeCun. 2026. World Models for Learning Dexterous Hand-Object Interactions from Human Videos. arXiv:2512.13644 <https://arxiv.org/abs/2512.13644v2> Accepted to ECCV 2026; arXiv version 2, revised 16 March 2026; official accepted-paper list pending publisher checks.
- [20] Jiarui Guan, Wenshuai Zhao, Yue Pei, Ziliang Chen, Arno Solin, and Juho Kannala. 2026. Point Tracking Improves World Action Models. arXiv:2605.23856 <https://arxiv.org/abs/2605.23856v1> Preprint, version 1, 22 May 2026.

- [21] Yanjiang Guo, Tony Lee, Lucy Xiaoyang Shi, Jianyu Chen, Percy Liang, and Chelsea Finn. 2026. VLA: Iterative Co-Improvement of Vision-Language-Action Policy and World Model. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*. PMLR, 37998–38011. <https://proceedings.mlr.press/v306/guo26i.html>
- [22] David Ha and Jürgen Schmidhuber. 2018. World Models. arXiv:1803.10122 <https://arxiv.org/abs/1803.10122>
- [23] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2020. Dream to Control: Learning Behaviors by Latent Imagination. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=S1lOTC4tDS>
- [24] Danijar Hafner, Timothy Lillicrap, Ian Fischer, Ruben Villegas, David Ha, Honglak Lee, and James Davidson. 2019. Learning Latent Dynamics for Planning from Pixels. In *Proceedings of the 36th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 97)*. PMLR, 2555–2565. <https://proceedings.mlr.press/v97/hafner19a.html>
- [25] Nicklas Hansen, Hao Su, and Xiaolong Wang. 2024. TD-MPC2: Scalable, Robust World Models for Continuous Control. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2024/hash/cf73d57b6dca32b293df7c2d5341f49-Abstract-Conference.html
- [26] Ryan Hoque, Ashwin Balakrishna, Ellen Novoseller, Albert Wilcox, Daniel S. Brown, and Ken Goldberg. 2022. ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning. In *Proceedings of the 5th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 164)*. PMLR, 598–608. <https://proceedings.mlr.press/v164/hoque22a.html>
- [27] Bohan Hou, Gen Li, Jindou Jia, et al. 2026. World Model for Robot Learning: A Comprehensive Survey. arXiv:2605.00080 <https://arxiv.org/abs/2605.00080v1>
- [28] Tianshuai Hu, Zeyang Gong, Lingdong Kong, Xiaodong Mei, Yiyi Ding, Qi Zeng, Ao Liang, Rong Li, Yangyi Zhong, and Junwei Liang. 2026. NavThinker: Action-Conditioned World Models for Coupled Prediction and Planning in Social Navigation. arXiv:2603.15359 <https://arxiv.org/abs/2603.15359v1> Preprint, version 1, 16 March 2026.
- [29] Bingqi Huang, Bingchuan Wei, Yingkai Cai, and Zhaokui Wang. 2026. Selective Cross-View Consistency for World Action Models: Held-Out Viewpoint Robustness Without Test-Time Camera Information. arXiv:2608.21402 [cs.RO] <https://arxiv.org/abs/2608.21402v1> Preprint, version 1, submitted 7 August 2026.
- [30] Suning Huang, Qianzhong Chen, Xiaohan Zhang, Jiankai Sun, and Mac Schwager. 2025. ParticleFormer: A 3D Point Cloud World Model for Multi-Object, Multi-Material Robotic Manipulation. In *Proceedings of the 9th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 305)*. PMLR, 4941–4957. <https://proceedings.mlr.press/v305/huang25c.html>
- [31] Weidong Huang, Jiaming Ji, Chunhe Xia, Borong Zhang, and Yaodong Yang. 2024. SafeDreamer: Safe Reinforcement Learning with World Models. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2024/hash/ece182f93af26c64187ba3f7dfd4309a-Abstract-Conference.html
- [32] Wen Huang, Haoran Sun, Yongjian Guo, Yunxuan Ma, Haoran Li, Jing Long, Zhouying Mo, Zhong Guan, Yucheng Guo, Shuai Di, and Junwu Xiong. 2026. NoiseGate: Learning Per-Latent Timestep Schedules as Information Gating in World Action Models. arXiv:2605.07794 [cs.RO] <https://arxiv.org/abs/2605.07794v1>
- [33] Zhefeng Huang, Yilin Cai, Ankit Patel, Mohammad Hajiha, Brendan Browne, and Yue Chen. 2026. Failure Detection for Surgical Robot Imitation Policies via Flow-Matching World Modeling. arXiv:2607.27511 <https://arxiv.org/abs/2607.27511v1> Preprint, version 1, 29 July 2026.
- [34] Michael Janner, Justin Fu, Marvin Zhang, and Sergey Levine. 2019. When to Trust Your Model: Model-Based Policy Optimization. In *Advances in Neural Information Processing Systems*, Vol. 32. <https://proceedings.neurips.cc/paper/2019/hash/5faf461eff3099671ad63cf3f094f7f-Abstract.html>
- [35] Gaytri Jena, Kapil Wanaskar, Vinija Jain, Aman Chadha, Vasu Sharma, and Amitava Das. 2026. Do World Models Make Better Robots? A Survey of Evaluation Benchmarks for Predictive Embodied Intelligence. arXiv:2609.29669 <https://arxiv.org/abs/2609.29669v1>
- [36] Fan Jia, Weixin Mao, Yingfei Liu, Yucheng Zhao, Yuqing Wen, Chi Zhang, Xiangyu Zhang, and Tiancai Wang. 2023. ADriver-I: A General World Model for Autonomous Driving. arXiv:2311.13549 [cs.CV] <https://arxiv.org/abs/2311.13549v1> Preprint, version 1, 22 November 2023.
- [37] Alexander Khazatsky, Karl Pertsch, Suraj Nair, et al. 2024. DROID: A Large-Scale In-The-Wild Robot Manipulation Dataset. In *Proceedings of Robotics: Science and Systems*. Delft, Netherlands. doi:10.15607/RSS.2024.XX.120
- [38] Byungjun Kim, Taeksoo Kim, Junyoung Lee, and Hanbyul Joo. 2026. Dexterous World Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 29663–29673. https://openaccess.thecvf.com/content/CVPR2026/html/Kim_Dexterous_World_Models_CVPR_2026_paper.html
- [39] Joonkyung Kim, Wenxi Chen, Davood Soleymanzadeh, Yi Ding, Xiangbo Gao, Zhengzhong Tu, Ruqi Zhang, Fan Fei, Sushant Veer, Yiwei Lyu, Minghui Zheng, and Yan Gu. 2026. Modular Safety Guardrails Are Necessary for Foundation-Model-Enabled Robots in the Real World. arXiv:2602.04056 <https://arxiv.org/abs/2602.04056v1> Preprint, version 1, 3 February 2026.
- [40] Moo Jin Kim, Yihui Gao, Tsung-Yi Lin, Yen-Chen Lin, Yunhao Ge, Grace Lam, Percy Liang, Shuran Song, Ming-Yu Liu, Chelsea Finn, and Jinwei Gu. 2026. Cosmos Policy: Fine-Tuning Video Models for Visuomotor Control and Planning. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2026/hash/748becc400a57c0e31cfe6a2e7951467-Abstract-Conference.html
- [41] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan P. Foster, Pannag R. Sanketi, Quan Vuong, Thomas Kollar, Benjamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn. 2025. OpenVLA: An Open-Source Vision-Language-Action Model. In *Proceedings of the 8th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 270)*. PMLR, 2679–2713. <https://proceedings.mlr.press/v270/kim25c.html>
- [42] Sven Kirchner, Nils Purschke, and Alois Knoll. 2026. A Survey of World Models for Physical AI with Uncertainty Representation and Control. *Discover Artificial Intelligence* 6, 1 (2026), 1037. doi:10.1007/s44163-026-02122-1

- [43] Geonmyeong Lee and Byoung-Tak Zhang. 2026. Beyond Task Success: Stage-Wise Reliability of World Model Planning under Sensing Degradation. arXiv:2609.07126 [cs.RO] <https://arxiv.org/abs/2609.07126v1> Preprint, version 1, submitted 7 September 2026.
- [44] Dayou Li, Jiuzhou Lei, Hao Wang, Lulin Liu, Yunhao Yang, Zihan Wang, Bangya Liu, Minghui Zheng, and Zhiwen Fan. 2026. Learning Actionable Manipulation Recovery via Counterfactual Failure Synthesis. arXiv:2603.13528 <https://arxiv.org/abs/2603.13528v1> Preprint, version 1, 13 March 2026.
- [45] Huanyu Li, Kun Lei, Sheng Zang, Kaizhe Hu, Yongyuan Liang, Bo An, Xiaoli Li, and Huazhe Xu. 2026. Failure-Aware RL: Reliable Offline-to-Online Reinforcement Learning with Self-Recovery for Real-World Manipulation. arXiv:2601.07821 <https://arxiv.org/abs/2601.07821v1> Preprint, version 1, 12 January 2026.
- [46] Runze Li, Hongyin Zhang, Junxi Jin, Qixin Zeng, Zifeng Zhuang, Yiqi Tang, Shangke Lyu, and Donglin Wang. 2026. World-Value-Action Model: Implicit Planning for Vision-Language-Action Systems. arXiv:2604.14732 <https://arxiv.org/abs/2604.14732v2> Preprint, version 2, 19 April 2026.
- [47] Shalfun Li, Victor Yao, Charles Yang, Truth Qu, Regis Cheng, Ryan Yu, Howard Lu, Newton Von, Vincent Chen, Yohann Tang, Maeve Zhang, Ellie Ma, Gody Li, Starrick Liu, Sage Yang, Lorient Shu, J. W. Gao, Ethan Chen, Colin Ye, Yu Sun, Elise Mon, PS Zhang, Neo Li, Lily Li, James Wang, Ping Yang, Chris Pan, Lucy Liang, Hang Su, Roy Gan, Hao Wang, and Qian Wang. 2026. WALL-WM: Carving World Action Modeling at the Event Joints. arXiv:2606.01955 [cs.RO] <https://arxiv.org/abs/2606.01955v2>
- [48] Yingyan Li, Shuyao Shang, Weisong Liu, Bing Zhan, Haochen Wang, Yuqi Wang, Yuntao Chen, Xiaoman Wang, Yasong An, Chufeng Tang, Lu Hou, Lue Fan, and Zhaoxiang Zhang. 2026. DriveVLA-W0: World Models Amplify Data Scaling Law in Autonomous Driving. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2026/hash/0d70423f59c5fdd24f0dd3fa52e34623-Abstract-Conference.html
- [49] Yaxuan Li, Zhongyi Zhou, Yefei Chen, Yanjiang Guo, Jiaming Liu, Shanghang Zhang, Jianyu Chen, and Yichen Zhu. 2026. Hi-WM: Human-in-the-World-Model for Scalable Robot Post-Training. arXiv:2604.21741 <https://arxiv.org/abs/2604.21741v2> Preprint, version 2.
- [50] Ao Liang, Lingdong Kong, Tianyi Yan, Hongsi Liu, Yu Yang, Ziqi Huang, Wei Yin, Jialong Zuo, Yixuan Hu, Dekai Zhu, Dongyue Lu, Youquan Liu, Guangfeng Jiang, Linfeng Li, Xiangtai Li, Long Zhuo, Lai Xing Ng, Benoit R. Cottreau, Changxin Gao, Liang Pan, Wei Tsang Ooi, and Ziwei Liu. 2026. WorldLens: Full-Spectrum Evaluations of Driving World Models in Real World. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. https://openaccess.thecvf.com/content/CVPR2026/papers/Liang_WorldLens_Full-Spectrum_Evaluations_of_Driving_World_Models_in_Real_World_CVPR_2026_paper.pdf
- [51] Haotian Liang, Xinyi Chen, Bin Wang, Mingkan Chen, Yitian Liu, Yuhao Zhang, Zhanxin Chen, Tianshuo Yang, Yilun Chen, Jiangmiao Pang, Dong Liu, Xiaokang Yang, Yao Mu, Wenqi Shao, and Ping Luo. 2026. MM-ACT: Learn from Multimodal Parallel Generation to Act. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 35080–35090. https://openaccess.thecvf.com/content/CVPR2026/html/Liang_MM-ACT_Learn_from_Multimodal_Parallel_Generation_to_Act_CVPR_2026_paper.html
- [52] Yihan Lin, Jiawei He, Shifeng Bao, Chen Zhao, Yang Li, Xiaobo Wang, Yan Wang, Cheng Chi, and Jing Zhang. 2026. JEPa-WAM: Learning Vision–Language–Action Policies with Joint-Embedding World Modeling. arXiv:2608.09381 [cs.RO] <https://arxiv.org/abs/2608.09381v1> Preprint, version 1, submitted 10 August 2026.
- [53] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. 2023. LIBERO: Benchmarking Knowledge Transfer for Lifelong Robot Learning. In *Advances in Neural Information Processing Systems*, Vol. 36. https://proceedings.neurips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html
- [54] Fei Liu, Shichao Xie, Minghua Luo, Zedong Chu, Junjun Hu, Xiaolong Wu, and Mu Xu. 2026. NavForesee: A Unified Vision-Language World Model for Hierarchical Planning and Dual-Horizon Navigation Prediction. arXiv:2512.01550 <https://arxiv.org/abs/2512.01550v2> Preprint, version 2, 13 March 2026.
- [55] Huihan Liu, Shivin Dass, Roberto Martín-Martín, and Yuke Zhu. 2024. Model-Based Runtime Monitoring with Interactive Imitation Learning. In *IEEE International Conference on Robotics and Automation*. IEEE, 4154–4161. doi:10.1109/ICRA57147.2024.10611038
- [56] Lin Liu, Zhicheng Bao, Lu Zhang, Ziyang Song, Wu Yang, Shuai Tao, Wulong Liu, and Huchuan Lu. 2026. LIBERO-RECOVER: Beyond Task Success Towards Failure Recovery in Robotic Manipulation Models. arXiv:2609.05178 <https://arxiv.org/abs/2609.05178v1> Preprint, version 1, 4 September 2026.
- [57] Xiaokang Liu, Zechen Bai, Hai Ci, Kevin Yuchen Ma, and Mike Zheng Shou. 2026. World-VLA-Loop: Closed-Loop Learning of Video World Model and VLA Policy. arXiv:2602.06508 <https://arxiv.org/abs/2602.06508v2> Preprint, version 2, 25 May 2026.
- [58] Hongbo Lu, Liang Yao, Chenghao He, Haoyu Wang, Xiang Gu, Xianfei Li, Wenlong Liao, Tao He, and Pai Peng. 2026. The DAWN of World-Action Interactive Models. arXiv:2605.11550 [cs.CV] <https://arxiv.org/abs/2605.11550v1>
- [59] Zuxing Lu, Hongjia Zhai, Guanzhi Wang, et al. 2026. World-Action Models for Robot Learning and Control: A Survey. arXiv:2609.16074 <https://arxiv.org/abs/2609.16074v1>
- [60] Kailang Ma, Heye Huang, Inhi Kim, and Kitae Jang. 2026. Rethinking World Models for Safety-Critical Embodied Systems. arXiv:2609.03774 <https://arxiv.org/abs/2609.03774v2> Perspective article, version 2, 1 October 2026.
- [61] Zipei Ma, Xiaofei Wei, Junzhe Jiang, Shunlin Lu, and Li Zhang. 2026. TacPAC: Tactile Prediction and Real-Time Action Correction in World-Action Models for Contact-Rich Manipulation. arXiv:2609.05266 [cs.RO] <https://arxiv.org/abs/2609.05266v1> Preprint, version 1, submitted 4 September 2026.
- [62] Oier Mees, Lukas Hermann, Erick Rosete-Beas, and Wolfram Burgard. 2022. CALVIN: A Benchmark for Language-Conditioned Policy Learning for Long-Horizon Robot Manipulation Tasks. *IEEE Robotics and Automation Letters* 7, 3 (2022), 7327–7334. doi:10.1109/LRA.2022.3180108

- [63] Shangchen Miao, Ningya Feng, Jialong Wu, Ye Lin, Xu He, Dong Li, and Mingsheng Long. 2026. JEPA-VLA: Video Predictive Embedding is Needed for VLA Models. arXiv:2602.11832 [cs.CV] <https://arxiv.org/abs/2602.11832v1> Preprint, version 1, submitted 12 February 2026.
- [64] Utkarsh A. Mishra, Yongxin Chen, Danfei Xu, Yang Liu, Xi Chen, and Jiayuan Mao. 2026. Understanding and Mitigating the Video-Action Generalization Gap via Temporal Ratio. arXiv:2607.08127 [cs.CV] <https://arxiv.org/abs/2607.08127v1> Preprint, version 1, submitted 9 July 2026.
- [65] Nana Yaw Owusu Ofori-Ampofo, Samira Ebrahimi Kahou, and Joseph Thekinen. 2026. REBOOT: From Failure to Recovery—A Dataset and Benchmark for Precision Assembly. arXiv:2609.22591 <https://arxiv.org/abs/2609.22591v1> Preprint, version 1, 18 September 2026.
- [66] Open X-Embodiment Collaboration et al. 2024. Open X-Embodiment: Robotic Learning Datasets and RT-X Models. In *IEEE International Conference on Robotics and Automation*. IEEE, 6892–6903. doi:10.1109/ICRA57147.2024.10611477
- [67] Yan Qin, Yue Chen, Wenwei Lin, Shujia Liu, Chuqiao Lyu, Kailun Su, Weiyang Jin, Chenze Yu, Ping Luo, Wenbo Ding, Tianxing Chen, and Renjing Xu. 2026. DexTouch-WM: Learning Action-Conditioned Tactile World Models from Human Touch for Dexterous Robot Manipulation. arXiv:2609.20649 <https://arxiv.org/abs/2609.20649v2> Preprint, version 2, 18 September 2026; IROS 2026 RoBoWoMo workshop lightning talk.
- [68] Zachary Ravichandran, Alexander Robey, Vijay Kumar, George J. Pappas, and Hamed Hassani. 2026. Safety Guardrails for LLM-Enabled Robots. *IEEE Robotics and Automation Letters* 11, 4 (2026), 4649–4656. doi:10.1109/LRA.2026.3667488
- [69] Allen Z. Ren, Anushri Dixit, Alexandra Bodrova, Sumeet Singh, Stephen Tu, Noah Brown, Peng Xu, Leila Takayama, Fei Xia, Jake Varley, Zhenjia Xu, Dorsa Sadigh, Andy Zeng, and Anirudha Majumdar. 2023. Robots That Ask For Help: Uncertainty Alignment for Large Language Model Planners. In *Proceedings of the 7th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 229)*. PMLR, 661–682. <https://proceedings.mlr.press/v229/ren23a.html>
- [70] Alvaro Sanchez-Gonzalez, Jonathan Godwin, Tobias Pfaff, Rex Ying, Jure Leskovec, and Peter Battaglia. 2020. Learning to Simulate Complex Physics with Graph Networks. In *Proceedings of the 37th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 119)*. PMLR, 8459–8468. <https://proceedings.mlr.press/v119/sanchez-gonzalez20a.html>
- [71] Yu Shang, Zhuohang Li, Yiding Ma, et al. 2026. WorldArena: A Unified Benchmark for Evaluating Perception and Functional Utility of Embodied World Models. arXiv:2602.08971 <https://arxiv.org/abs/2602.08971>
- [72] Yu Shang, Yinzhou Tang, Yiding Ma, et al. 2026. WorldArena 2.0: Extending Embodied World Model Benchmarking on Modality, Functionality and Platform. arXiv:2605.17912 <https://arxiv.org/abs/2605.17912v1>
- [73] Stuart Shanks, Jonathan Embley-Riches, Jianheng Liu, Andromachi Maria Delfaki, Carlo Ciliberto, and Dimitrios Kanoulas. 2025. DreamerNav: Learning-Based Autonomous Navigation in Dynamic Indoor Environments Using World Models. *Frontiers in Robotics and AI* 12 (2025), 1655171. doi:10.3389/frobt.2025.1655171
- [74] Chen Shi, Jinrui Xu, Shaoshuai Shi, Kehua Sheng, Bo Zhang, and Li Jiang. 2026. DriveWAM: Video Generative Priors Enable Scalable World-Action Modeling for Autonomous Driving. arXiv:2605.28544 [cs.CV] <https://arxiv.org/abs/2605.28544v1>
- [75] Haisheng Su, Zongdai Liu, Xin Jin, Haoxuan Dou, Chengming Hu, Baorun Li, Zhanwang Liu, Ruiyan Xu, Jianjie Fang, Xin Zhang, Zhenjie Yang, Xue Yang, Chen Gao, Junchi Yan, Yong Li, and Wei Wu. 2026. WorldScape Policy 2.0: Empowering Steerable World Action Modeling with Reasoning-Augmented Memory. arXiv:2607.18840 <https://arxiv.org/abs/2607.18840v1> Preprint, version 1, 20 July 2026.
- [76] Jiayi Su, Yixin Zheng, Mi Yan, Li Yi, Zhizheng Zhang, and He Wang. 2026. GPT 6 Astra as an Embodied Policy. Technical report and code. <https://github.com/anonymous-report-421/GPT-as-Policy> Accessed 1 October 2026.
- [77] Jingwen Sun, Wenyao Zhang, Zekun Qi, Shaojie Ren, Zezhi Liu, Hanxin Zhu, Guangzhong Sun, Xin Jin, and Zhibo Chen. 2026. VLA-JEPA: Enhancing Vision–Language–Action Model with Latent World Model. arXiv:2602.10098 [cs.RO] <https://arxiv.org/abs/2602.10098v2> Preprint, version 2, revised 14 February 2026; the arXiv HTML title page displays 24 August 2026.
- [78] Yu Sun, Meng Cao, Yang Ping, Kaidong Zhang, Qingxuan Chen, Rongtao Xu, Liangwang Ruan, Xuecheng Chen, Dongxiu Liu, Yunxiao Yan, Zunnan Xu, Runze Xu, Charles Yang, Peilun Zhang, Xiaofan Li, Ruyi Gan, Liang Ma, Yuehao Yin, Jincheng Yu, Lufang Chen, Yuxin Liang, Peng Zhai, Hao Wang, Ivan Laptev, Ian Reid, Qian Wang, and Xiaodan Liang. 2026. ManipArena: Comprehensive Real-world Evaluation of Reasoning-Oriented Generalist Robot Manipulation. arXiv:2603.28545 [cs.RO] <https://arxiv.org/abs/2603.28545v2> Preprint, version 2, revised 1 July 2026.
- [79] Brijen Thananjeyan, Ashwin Balakrishna, Suraj Nair, Michael Luo, Krishnan Srinivasan, Minh Hwang, Joseph E. Gonzalez, Julian Ibarz, Chelsea Finn, and Ken Goldberg. 2021. Recovery RL: Safe Reinforcement Learning with Learned Recovery Zones. *IEEE Robotics and Automation Letters* 6, 3 (2021), 4915–4922. doi:10.1109/LRA.2021.3070252
- [80] Garrett Thomas, Yuping Luo, and Tengyu Ma. 2021. Safe Reinforcement Learning by Imagining the Near Future. In *Advances in Neural Information Processing Systems*, Vol. 34. <https://proceedings.neurips.cc/paper/2021/hash/73b277c11266681122132d024f53a75b-Abstract.html>
- [81] Rahul Venkatesh, Klemen Kotar, Lilian Naing Chen, Wanhee Lee, Gia Ancone, Seungwoo Kim, Luca Thomas Wheeler, Jared Watrous, Honglin Chen, Daniel Bear, Stefan Stojanov, and Daniel LK Yamins. 2026. Physical Object Understanding with a Physically Controllable World Model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2593–2602. https://openaccess.thecvf.com/content/CVPR2026/html/Venkatesh_Physical_Object_Understanding_with_a_Physically_Controllable_World_Model_CVPR_2026_paper.html
- [82] Alex Wang, Zhiwei Dong, Qicheng Bai, Chenshi Zhang, Yujie Yi, Guang Dai, Yong Liu, and Mengmeng Wang. 2026. DynBridge: Bridging Imagination and Control through Interaction Dynamics for Robot Manipulation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 22423–22432. https://openaccess.thecvf.com/content/CVPR2026/html/Wang_DynBridge_Bridging_Imagination_and_Control_through_Interaction_Dynamics_for_Robot_CVPR_2026_paper.html
- [83] Jiaxu Wang, Yicheng Jiang, Tianlun He, Jingkai Sun, Qiang Zhang, Jiahang Cao, Zesen Gan, Mingyuan Sun, Qiming Shao, and Xiangyu Yue. 2026. MVISTA-4D: View-Consistent 4D World Model with Test-Time Action Inference for Robotic Manipulation. In *Proceedings of the 43rd International*

- Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306). PMLR, 131846–131866. <https://proceedings.mlr.press/v306/wang26jw.html>
- [84] Jun Wang, Yevgeniy Vorobeychik, and Yiannis Kantaros. 2026. CoFineLLM: Conformal Finetuning of LLMs for Language-Instructed Robot Planning. In *Proceedings of the 8th Annual Learning for Dynamics and Control Conference (Proceedings of Machine Learning Research, Vol. 331)*. PMLR, 1558–1574. <https://proceedings.mlr.press/v331/wang26c.html>
- [85] Yuqi Wang, Jiawei He, Lue Fan, Hongxin Li, Yuntao Chen, and Zhaoxiang Zhang. 2024. Driving into the Future: Multiview Visual Forecasting and Planning with World Model for Autonomous Driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 14749–14759. https://openaccess.thecvf.com/content/CVPR2024/html/Wang_Driving_into_the_Future_Multiview_Visual_Forecasting_and_Planning_with_CVPR_2024_paper.html
- [86] Yixuan Wang, Rhythm Syed, Fangyu Wu, Mengchao Zhang, Aykut Onol, Jose Barreiros, Hooshang Nayyeri, Tony Dear, Huan Zhang, and Yunzhu Li. 2026. Interactive World Simulator for Robot Policy Training and Evaluation. arXiv:2603.08546 <https://arxiv.org/abs/2603.08546v1> Preprint, version 1, 9 March 2026.
- [87] Zixing Wang, Kausik Sivakumar, Jinghuan Shang, Yafei Hu, Zhaoming Xie, Ran Gong, Xiaohan Zhang, and Karl Schmeckpeper. 2026. Efficient Sim-to-Real Transfer of World-Action Models from Synthetic Priors. arXiv:2606.31101 [cs.RO] <https://arxiv.org/abs/2606.31101v1> CVPR 2026 Embodied AI Workshop early result; arXiv version 1, submitted 30 June 2026.
- [88] Philipp Wu, Alejandro Escontrela, Danijar Hafner, Pieter Abbeel, and Ken Goldberg. 2023. DayDreamer: World Models for Physical Robot Learning. In *Proceedings of the 6th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 205)*. PMLR, 2226–2240. <https://proceedings.mlr.press/v205/wu23c.html>
- [89] Sheng Xu, Ruixing Jin, Huayi Zhou, Bo Yue, Guanren Qiao, Yueci Deng, Yunxin Tai, Kui Jia, and Guiliang Liu. 2026. From Reaction to Anticipation: Proactive Failure Recovery through Agentic Task Graph for Robotic Manipulation. arXiv:2605.11951 <https://arxiv.org/abs/2605.11951v1> Preprint, version 1, 12 May 2026.
- [90] Yuquan Xue, Le Xu, Zeyi Liu, Zhenyu Wu, Zhengyi Gu, Xinyang Song, Bofang Jia, and Ziwei Wang. 2026. WorldSample: Closed-loop Real-robot RL with World Modelling. arXiv:2607.02431 <https://arxiv.org/abs/2607.02431v1> Preprint, version 1, 2 July 2026.
- [91] Bowen Yan, Jiahao Xiao, Kehui Liu, Jianbo Zhang, Zicheng Zhang, Qi Jia, Zhongjie Jia, Haoming Song, Chunyi Li, Bin Zhao, and Guangtao Zhai. 2026. Can VLMs Diagnose and Recover from VLA Manipulation Faults?. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*. PMLR, 143967–143983. <https://proceedings.mlr.press/v306/yan26r.html>
- [92] Sherry Yang, Yilun Du, Seyed Ghasemipour, Jonathan Tompson, Leslie Kaelbling, Dale Schuurmans, and Pieter Abbeel. 2024. Learning Interactive Real-World Simulators. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2024/hash/c4d66eae503694424123b93ac0fbaf17-Abstract-Conference.html
- [93] Yu Yang, Jianbiao Mei, Yukai Ma, Siliang Du, Wenqing Chen, Yijie Qian, Yuxiang Feng, and Yong Liu. 2025. Driving in the Occupancy World: Vision-Centric 4D Occupancy Forecasting and Planning via World Models for Autonomous Driving. *Proceedings of the AAAI Conference on Artificial Intelligence* 39, 9 (2025), 9327–9335. doi:10.1609/aaai.v39i9.33010
- [94] Seonghyeon Ye, Yunhao Ge, Kaiyuan Zheng, et al. 2026. World Action Models are Zero-shot Policies. arXiv:2602.15922 <https://arxiv.org/abs/2602.15922v1>
- [95] Jiaju Yin, Zhenhui Zhang, Lixin Xu, Heng Zhang, Jun Shao, Yating Feng, Arash Ajoudani, and Renjing Xu. 2026. How to Learn from What a Human Would Avoid? Intervention-Aware World Models with Real-World RL for Dexterous Manipulation. arXiv:2609.06009 [cs.RO] <https://arxiv.org/abs/2609.06009v1> Accepted to CoRL 2026; arXiv version 1, submitted 5 September 2026.
- [96] Dongjie Yu, Wenjun Zou, Yujie Yang, Haitong Ma, Shengbo Eben Li, Yuming Yin, Jianyu Chen, and Jingliang Duan. 2024. Safe Model-Based Reinforcement Learning With an Uncertainty-Aware Reachability Certificate. *IEEE Transactions on Automation Science and Engineering* 21, 3 (2024), 4129–4142. doi:10.1109/TASE.2023.3292388
- [97] Haoran Yuan, Weigang Yi, Zhenyu Zhang, Wendi Chen, Yuchen Mo, Jiashi Yin, Xinzhuo Li, Xiangyu Zeng, Chuan Wen, Cewu Lu, Katherine Driggs-Campbell, and Ismini Lourentzou. 2026. VTAM: Video-Tactile-Action Models for Complex Physical Interaction Beyond VLAs. arXiv:2603.23481 <https://arxiv.org/abs/2603.23481v1> Preprint, version 1, 24 March 2026.
- [98] Xianchao Zeng, Xinyu Zhou, Yucheng Li, Jiayou Shi, Tianle Li, Liangming Chen, Lei Ren, and Yong-Lu Li. 2026. Diagnose, Correct, and Learn from Manipulation Failures via Visual Symbols. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 42386–42395. https://openaccess.thecvf.com/content/CVPR2026/html/Zeng_Diagnose_Correct_and_Learn_from_Manipulation_Failures_via_Visual_Symbols_CVPR_2026_paper.html
- [99] Fan Zhang and Michael Gienger. 2026. Learning Robot Manipulation from Audio World Models. arXiv:2512.08405 <https://arxiv.org/abs/2512.08405v1> Accepted to an ICRA 2026 workshop on multimodal manipulation; arXiv version 1, 9 December 2025; no DOI listed in the author-hosted accepted manuscript.
- [100] Haoran Zhang, Yifu Lu, Boyang Wang, Xuhui Kang, Yen-Ling Kuo, Zezhou Cheng, Mengdi Wang, and Odest Chadwicke Jenkins. 2026. Foresight: Failure Detection for Long-Horizon Robotic Manipulation with Action-Conditioned World Model Latents. arXiv:2606.23085 <https://arxiv.org/abs/2606.23085v1> Preprint, version 1, 22 June 2026.
- [101] Wenbo Zhang, Kaixuan Wang, Yutao Ouyang, et al. 2026. An Unexpected Robot Policy: Early Evaluations of GPT-6 Astra on RoboDojo and Beyond. arXiv:2609.24170 <https://arxiv.org/abs/2609.24170v1> Version 1, 21 September 2026.

- [102] Yang Zhang, Jiangyuan Zhao, Chenyou Fan, Jiayu Hu, Xiu Yuan, Chenjia Bai, and Xiu Li. 2026. V-JEPA Policy: Building Effective World–Action Models on Predictive Visual Latents. arXiv:2609.37250 [cs.CV] <https://arxiv.org/abs/2609.37250v1> Preprint, version 1, submitted 29 September 2026.
- [103] Zeyu Zhang, Danning Li, Ian Reid, and Richard Hartley. 2026. GeoWorld: Geometric World Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 30952–30963. https://openaccess.thecvf.com/content/CVPR2026/html/Zhang_GeoWorld_Geometric_World_Models_CVPR_2026_paper.html
- [104] Zhanguang Zhang, Zhiyuan Li, Behnam Rahmati, Rui Heng Yang, Yintao Ma, Amir Rasouli, Sajjad Pakdamansavoji, Yangzheng Wu, Lingfeng Zhang, Tongtong Cao, Feng Wen, Xinyu Wang, Xingyue Quan, and Yingxue Zhang. 2026. Do World Action Models Generalize Better than VLAs? A Robustness Study. arXiv:2603.22078 [cs.RO] <https://arxiv.org/abs/2603.22078v5> Preprint, version 5, revised 30 July 2026.
- [105] Zhiyuan Zhang, Pokuang Zhou, Kaidi Zhang, Adeesh Desai, Temitope Amosa, Davood Soleymanzadeh, Jiuzhou Lei, Yuhao Zhou, Minghui Zheng, and Yu She. 2026. ContactWorld: What Representations Matter for Vision–Tactile Latent World Models in Contact-Rich Manipulation. arXiv:2606.13877 <https://arxiv.org/abs/2606.13877v3> Preprint, version 3, 24 September 2026.
- [106] Zesen Zhao, Minkyung Cho, Hui Shen, Boyuan Zheng, Kunxiao Gao, Yulong Cao, and Z. Morley Mao. 2026. Test-Time Scaling for World Action Models via Zero-Shot Geometric Evaluation. arXiv:2607.17454 [cs.RO] <https://arxiv.org/abs/2607.17454v1> Preprint, version 1; HTML and PDF use “Geometric Verification” in the title.
- [107] Haoyu Zhen, Qiao Sun, Hongxin Zhang, Junyan Li, Siyuan Zhou, Yilun Du, and Chuang Gan. 2025. Learning 4D Embodied World Models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*. 5337–5347. https://openaccess.thecvf.com/content/ICCV2025/html/Zhen_Learning_4D_Embodied_World_Models_ICCV_2025_paper.html
- [108] Wenzhao Zheng, Weiliang Chen, Yuanhui Huang, Borui Zhang, Yueqi Duan, and Jiwen Lu. 2024. OccWorld: Learning a 3D Occupancy World Model for Autonomous Driving. In *European Conference on Computer Vision*. 55–72. https://www.ecva.net/papers/eccv_2024/papers_ECCV/html/2024_ECCV_2024_paper.php
- [109] Yuhang Zheng, Songen Gu, Yupeng Zheng, Weize Li, Yujie Zang, Shuai Tian, Xiang Li, Ce Hao, Chen Gao, Si Liu, Haoran Li, Yilun Chen, Shuicheng Yan, and Wenchao Ding. 2026. OmniVTA: Visuo-Tactile World Modeling for Contact-Rich Robotic Manipulation. arXiv:2603.19201 <https://arxiv.org/abs/2603.19201v3> Preprint, version 3, 10 August 2026.
- [110] Zirui Zheng, Jiaqian Yu, Xiongfeng Peng, Jun Shi, Mingyi Li, Chao Zhang, Weiming Li, Dong Wang, Huchuan Lu, and Xu Jia. 2026. Mem-World: Memory-Augmented Action-Conditioned World Models for Persistent Robot Manipulation. arXiv:2606.18960 <https://arxiv.org/abs/2606.18960v1> Preprint, version 1, 17 June 2026.
- [111] Jianyi Zhou, Feiyang Hong, Yunhao Li, Yicheng Zhao, Yongjue Cen, Zirui Liu, Jiakang Huang, Zirui Chen, Ruiyang Zhang, Weizhuo Zhu, Xuhua Song, and Shuo Yang. 2026. TouchWorld: A Predictive and Reactive Tactile Foundation Model for Dexterous Manipulation. arXiv:2607.07287 <https://arxiv.org/abs/2607.07287v2> Preprint, version 2, 9 July 2026.
- [112] Jiaming Zhou, Qihang Zhang, Gangwei Xu, Cunxin Fan, Yujie Zhao, Ruilin Wang, Yiming Luo, Shuai Yang, Xing Zhu, Yujun Shen, Junwei Liang, and Yinghao Xu. 2026. Zero-WAM: In-Context World-Action Modeling from Human Videos for Open-Ended Task Generalization. arXiv:2608.26103 [cs.RO] <https://arxiv.org/abs/2608.26103v2> Preprint, version 2, revised 27 August 2026.
- [113] Siyuan Zhou, Hejun Wang, Hu Cheng, Jinxi Li, Dongsheng Wang, Junwei Jiang, Yixiao Jin, Jiayue Huang, Shiwei Mao, Shangjia Liu, Yafei Yang, Hongkang Song, Shenxing Wei, Zihui Zhang, Bing Wang, Zhihua Wang, Chuhang Zou, and Bo Yang. 2026. PhysInOne: Visual Physics Learning and Reasoning in One Suite. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 33131–33142. https://openaccess.thecvf.com/content/CVPR2026/html/Zhou_PhysInOne_Visual_Physics_Learning_and_Reasoning_in_One_Suite_CVPR_2026_paper.html
- [114] Yang Zhou, Xiaofeng Wang, Hao Shao, Letian Wang, Guosheng Zhao, Jiangnan Shao, Jiagang Zhu, Tingdong Yu, Zheng Zhu, Guan Huang, and Steven L. Waslander. 2026. DriveDreamer-Policy: A Geometry-Grounded World–Action Model for Unified Generation and Planning. arXiv:2604.01765 <https://arxiv.org/abs/2604.01765v1> Preprint, version 1, 2 April 2026.